Do the poor pay more for increasing market concentration? A study of retail petroleum

Peter Ormosi & Franco Mariuzzo
Centre for Competition Policy and School of Law,
University of East Anglia

CCP Working Paper 21-08
This version: 30 June 2021

One of the central tenets of industrial organisation is that increasing/decreasing market concentration is likely to lead to increased/reduced markups. But does this affect every consumer to the same extent? Previous literature agrees that there can be significant price dispersion even in the case of homogeneous goods, which is at least partially due to the heterogeneity in how much consumers engage with the market. We link this heterogeneity to the impact of changing market concentration on markups. For this purpose, we employ a combination of 18 years of station-level motor fuel price data from Western Australia and a rich set of information on local market concentration. We summon a non-parametric causal forest approach to explore the heterogeneity in the effect of market exit/entry. The paper offers evidence of the distributional effect of changing market concentration. Areas with lower income experience a larger increase in petrol stations' price margin as a result of market exit. On the other hand, entry does not benefit the same low-income areas with a larger reduction in the margin than in high-income areas. We argue that these findings are due to differences in how much consumers in different demographic groups engage with the market. Our findings give support to the argument that antitrust could help address inequality while staying true to its mission of promoting competition, provided that priorities are given to not only fixing supply-side problems but also to exploring demand-side remedies..

Contact Details:

Peter Ormosi	P.Ormosi@uea.ac.uk
Franco Mariuzzo	F.Mariuzzo@uea.ac.uk

ISSN 1745-9648

Do the poor pay more for increasing market concentration? A study of retail petroleum markets*

Franco Mariuzzo
School of Economics
University of East Anglia
NR4 7TJ, Norwich, UK
f.mariuzzo@uea.ac.uk

Peter L Ormosi[†]
Norwich Business School
University of East Anglia
NR4 7TJ, Norwich, UK
p.ormosi@uea.ac.uk

June 30, 2021

Abstract

One of the central tenets of industrial organisation is that increasing/decreasing market concentration is likely to lead to increased/reduced markups. But does this affect every consumer to the same extent? Previous literature agrees that there can be significant price dispersion even in the case of homogeneous goods, which is at least partially due to the heterogeneity in how much consumers engage with the market. We link this heterogeneity to the impact of changing market concentration on markups. For this purpose, we employ a combination of 18 years of station-level motor fuel price data from Western Australia and a rich set of information on local market concentration. We summon a non-parametric causal forest approach to explore the heterogeneity in the effect of market exit/entry. The paper offers evidence of the distributional effect of changing market concentration. Areas with lower income experience a larger increase in petrol stations' price margin as a result of market exit. On the other hand, entry does not benefit the same low-income areas with a larger reduction in the margin than in high-income areas. We argue that these findings are due to differences in how much consumers in different demographic groups engage with the market. Our findings give support to the argument that antitrust could help address inequality while staying true to its mission of promoting competition, provided that priorities are given to not only fixing supply-side problems but also to exploring demand-side remedies.

Keywords inequality · market concentration · income · consumer search · causal forests

1 Introduction

Among the most intensively researched theoretical and empirical questions in the area of industrial organisation is the relationship between market concentration and prices. There is ample evidence showing that increasing concentration leads to increased prices and more competition lowers prices. This has also been linked more directly to market exit and entry. These findings have served as the basis for economic policies around the world, to liberalise markets and promote competition to improve the efficiency of the economy, and ultimately to maximise the welfare of consumers.

In this paper, our interest shifts away from the average consumer, who has been the central subject of these previous works. Instead, we ask the question: does increasing or falling market concentration affect every consumer the same way? To answer this question, we look at a homogeneous good, petroleum. Walrasian theory would suggest that in a homogeneous goods market, consumers pay the same price, therefore if entry and exit affect prices, the price change will be the same for every consumer. But extensive search literature has proven that this is not the case if consumers

*The paper has benefited from comments from participants of the CCP Seminar Series (2020), and the University of Amsterdam conference (2021), Should Wealth and Income Inequality Be a Competition Law Concern? All data and code used in this paper will be given open access upon publication of the paper.

†Corresponding author.

differ in how much they engage with the market (Salop & Stiglitz 1977, Varian 1980), and in their willingness to pay (Diamond 1987). If finding low prices requires consumers to engage in costly search, and if consumers differ in their willingness to pay, economic theory and evidence show that in equilibrium it can be optimal for different stores to charge different prices for the same good. Numerous empirical works have found support for these arguments (Woodward & Hall 2012, Allen et al. 2014, Lach & Moraga-González 2017, Wilson & Price 2010, Stango & Zinman 2016). A stylised and somewhat simplistic synopsis of these papers would be that a high willingness to search is associated with lower prices and vice versa. Logically, it would follow from these findings that if changing market concentration changes the equilibrium price, the price change will reflect the heterogeneity in consumers' engagement with the market.

Our main research question is to find out whether this is the case, and more specifically, whether heterogeneity in consumer engagement leads to distributional effects, i.e. do lower-income households pay more or less for increasing concentration, and do they benefit more or less from an increase in competition? This would depend on whether low/high-income households have low or high willingness to pay and whether they are more or less likely to search. The answer is not intuitively obvious. For example, households on higher income may be expected to search more, and pay less than poor households, or the other way around: richer households have higher consumption and so stand to make higher absolute savings from search, but their opportunity costs of time are likely to be higher than poorer households. Earlier works that have linked search and income - reviewed by Byrne & Martin (2021) - offer somewhat mixed evidence, although they point more towards the conclusion that lower-income consumers search less.

To provide novel empirical evidence to this question, we explore the geographical and over-time variation of market structure and prices in local petrol retail markets in Western Australia to design a natural experiment to study the impact of increasing and decreasing local market concentration on the retail margin. In this process we make use of the variation in local demographic characteristics, to investigate the relationship between the impact of changing market concentration and demand-side heterogeneity (with a particular focus on income). By setting up our study as an event study design, we can adjust the event window in a way that ensures that exit and entry can be considered exogenous in our experiments. This employs the assumption that exit and entry are not driven by short-term changes in the retail margin, something that we support with our data, and therefore shortening the event window reduces the risk of reverse causality biasing our findings.

Our paper ties up with and makes contributions to four main strands of work. Firstly, we offer new evidence to the growing body of literature on the distributional impact of market power. One of the pioneering papers on this, Baker & Salop (2015) set out the problem and offered an agenda for further work. Some of these works focus on the link between market power and inequality (Ennis et al. 2019, Khan & Vaheesan 2017). Despite the increased focus, much of the currently available analysis is either purely conceptual or is based on empirical work conducted on aggregated macro data (Ennis et al. 2019, Zac et al. 2020, Dierx et al. 2017). There is much less evidence at the market level, which is hardly surprising; linking changes in market concentration to different demographic groups is not a trivial exercise given its intensive data requirements. We approach this problem differently, motivated by the question: do the poor pay more for things, the leading thought behind the path-breaking book by Caplovitz (1963). By investigating some of the sources that would explain why some people pay a poverty premium, we ask whether the poor pay more if market concentration drives up prices. We view the main contribution of this paper is the rare, market-level evidence, on how local variation in the effect of market exit and entry can be linked to local variation in income.

Second, we draw from the literature on price dispersion in homogeneous goods. These works are motivated by the observation that price dispersion exists, even with homogeneous goods and where price information is readily available. Close to our work is Pennerstorfer et al. (2020), who formulate a theoretical model in which they link consumer information to price dispersion and predict an inverted U-shaped relationship. They test this model with retail petrol price data about Austria. To proxy for consumer information, they assume that drivers with longer commuting distances are more informed of petrol prices than those who commute less. To look at the impact of changing market concentration, Lach & Moraga-González (2017) used data from the Dutch petrol retail market. They find that price distributions in less competitive markets first-order stochastically dominate price distributions in more competitive markets and that consumer gains from the increasing competition are larger for more informed consumers. Byrne & de Roos (2017) look at the intensity of search in the petrol market, using the same data as ours - but accessing the number of visits to the FuelWatch website and instead of changes in market structure, investigate search intensity as a function of price dispersion. To add to this stream of literature, instead of focusing on the magnitude of price dispersion, we look at the impact of changes in market concentration on the expected price (profit margin). Our data allows us to account for numerous factors that drive the heterogeneity in the retail price effect of changing market concentration.

Thirdly, we link with works on the price impact of market exit and entry. Regarding a specific type of exit (merger), reviews of the large body of ex-post mergers studies, such as Kwoka (2014), or Mariuzzo & Ormosi (2019) reveal the depth of evidence establishing the link between mergers and price changes. Of the studies that looked at retail petroleum markets, the evidence is mixed. Simpson & Taylor (2008) find no evidence of higher prices, Hastings (2004) and Taylor et al. (2010) find price increases of different magnitudes. Regarding entry, Barron et al. (2004) show that adding one fuel station within a local market (i.e., 2.4 km ring) leads to a price reduction that varies across cities from 1.84 to 5.26 US cents per litre. However, Hosken et al. (2008) use a larger dataset and find no relationship between firm density and market price. All of these works study the impact of mergers or market entry on the average consumer. Allen et al. (2014) argues that this average effect of mergers (in their application to mortgage markets) underestimates the increase in market power, and they show that competition benefits only consumers at the bottom and middle of the transaction price distribution. We posit that the approach we use in this paper could be employed (and frequently the data is already available) in many merger retrospective studies to gain an understanding of the distributional impact of these transactions.

Finally, we make a methodological contribution to the literature on using causal forests to estimate heterogeneous treatment effects, as proposed by Athey et al. (2019). Whereas a large number of location, firm, and time characteristics in our data raises dimensionality issues, which would justify the use of a tree-based approach, at the same time, we have a relatively small cross-section of exits and entries in our sample. To handle this problem, we propose using an ensemble causal forest approach. We demonstrate through simulations, that this performs better than a single causal forest in cases with low number of observations and large number of estimable parameters.

Looking at petroleum retail margins for two products (unleaded petrol, and diesel) in Western Australia, between 2001 and 2019, we find that, in line with conventional industrial organisation theory, exit leads to an increase (although not significant on the average), and entry causes a drop in the retail margin. When dissecting our estimates to explore the heterogeneity in these findings, we find that low-income households experience a larger (and significant) increase in the price margin with exit. At the same time, they do not benefit from a lower drop in the margin with entry. This stands to suggest that low-income households either have a higher willingness to pay, or lower willingness to search, or both. We also find that concentrated markets witness a larger rise/fall in the price margin from exit/entry. Other factors, such as commuting distance, age, education also drive some of the heterogeneity but even after controlling for these factors, the difference between low and high-income households remains.

These findings offer evidence that the current antitrust toolkit is adequate to help reduce inequalities, but a reconsideration of some of the conventional thinking around competition policies may be warranted. The lack of engagement of lower-income consumers with the market means that the deterrent effect of competition policy (since it deters increasing market concentration) can help slow increasing inequalities. Rigorous merger control that stops an increase in concentration would have the same effect. But restoring the level of competition (for example through enforcement action, or by breaking up monopolies) will only reduce inequalities if paired with adequate demand-side remedies.

The paper is organised as follows. First, we provide a stylised economic framework, which pulls together some of the canonical theories from previous literature. This is followed by an introduction and description of our data, and a discussion of the methodology. We then present the results of our causal forest estimates before offering results from a linear regression, which also allows us to offer results that account for the potential endogeneity of exit and entry.

2 The economic framework

To understand why increasing or falling market concentration would have a distributional effect, we draw on Pennerstorfer et al. (2020), who builds on Varian (1980) and Stahl (1989) to model search in homogeneous goods, where consumers differ in their degree of informedness, and extend this model where consumers differ also in their willingness to pay for the product.[3] This allows distinction between consumers based on how informed they are about the price of petrol. For some, obtaining an additional price quote is costly; others are aware of all prices charged in the relevant market as they have access to the "clearinghouse". This setup leads to a mixed equilibrium due to the tension between charging a high price to exploit uninformed consumers and charging a low price to attract informed consumers. As a result, the authors find an inverse-U shape relationship between price dispersion and the proportion of informed consumers. Unlike Pennerstorfer et al. (2020), our focus is on the change in price (markup) as a result of a change in the number of suppliers in the market, rather than the change in price dispersion. Nevertheless, their model offers an appropriate architecture for us to express the expected price as a function of the number of stores and the proportion of informed consumers.

[3]It may seem obvious to link our work to works on the demand elasticity of petrol, such as Wadud et al. (2010), who provide estimates of motor fuel elasticities for different income levels. They look at the heterogeneity in petroleum demand elasticity, and find, among others an inverse relationship between income and demand elasticity. Whilst this bears some relevance to our study, it is tangential to our research question, as we are interested in how consumers choose between different suppliers, i.e. the brand elasticity of petrol demand, rather than product elasticity.

2.1 Informed and uninformed homogeneous consumers

In a market, there are N symmetric stores selling a homogeneous product. There is a mass of consumers, normalised to one, with a unit demand for the product. Initially, assume consumers to be homogeneous, all having the willingness to pay price v for the product. A proportion $\mu \in (0,1)$ of consumers are informed and therefore observe all prices set by the N stores in the market and buy at the lowest price if this price does not exceed their willingness to pay for the product. The uninformed consumers observe the price of one of the gas stations and to gain access to the price of any other gas station, they will have to search sequentially, facing the search cost s for any new price. The search cost is assumed to be a diminishing function of the number of stores in the market.[4] This means that the more stores there are, the closer a store gets to the consumer. In a sequential search, a consumer can optimise the distance (and time spent) between the stores they search.

The mixed strategy symmetric equilibrium price distribution $F(p)$ is the function that equates the profits achieved when a store charges a price in the interval $\left[\underline{p}, \overline{p}\right]$ to all consumers, and the profits of a store choosing to serve $\frac{1}{N}$ of the uninformed consumers at the high price $\overline{p}$,

$$(p-c)\left(\mu\left(1-F\left(p\right)\right)^{N-1}+\left(1-\mu\right)\frac{1}{N}\right)=\left(\overline{p}-c\right)\left(1-\mu\right)\frac{1}{N}, \tag{1}$$

which is,

$$F(p)=1-\left(\frac{1-\mu}{\mu}\frac{1}{N}\frac{\overline{p}-p}{p-c}\right)^{\frac{1}{N-1}}. \tag{2}$$

Invertibility of $F\left(p\right)$ along with the condition $F\left(\underline{p}\right)=0$ are sufficient to derive the lower bound for the price $\underline{p}=c+\frac{\overline{p}-c}{1+\frac{\mu}{1-\mu}N}$. The upper bound of the price support depends on the optimal reservation price, r, of the uninformed consumers, i.e., the price that makes these consumers indifferent between buying at r or incurring the search cost, s, to access a price drawn from the price distribution $F(p)$. This is captured by the following equation,

$$v-r=v-s(N)-\int_{\underline{p}}^{r}dF(p)-\left(1-F\left(r\right)\right)r. \tag{3}$$

As explained in Pennerstorfer et al. (2020), the reservation price that solves equation (3) is $r=c+\frac{s}{1-A}$, with $A=\int_0^1\frac{1}{1+\frac{\mu}{1-\mu}Nz^{N-1}}dz\in(0,1)$.

In this setting, the resulting expected price in the market is:

$$E(p)=\int_{\underline{p}}^{\min(r,v)}pdF(p)=c+\left(\min\left(r,v\right)-c\right)A. \tag{4}$$

The expected price increases monotonically with the proportion of uninformed consumers μ, $\frac{\partial E(p)}{\partial\mu}>0$. Under the assumption of the search cost being negatively related to the number of competitors, the relationship between price and the number of competitors can be negative. A necessary condition for this being the case is that the number of stores' elasticity of the reservation price (in absolute value), $\frac{\partial r}{\partial N}\frac{N}{r}$, is sufficiently large, i.e. it exceeds this other elasticity $\frac{\partial A}{\partial N}\frac{N}{A}$. This would mean that the relationship between search and cost produces $\frac{\partial E(p)}{\partial N}<0$.[5] This gives us our first testable hypothesis.

Hypothesis 1. *Exit ($dN<0$) has an incremental effect on the price. Entry ($dN>0$) has a regressive effect on the price.*

[4] This is not a far-fetched assumption. With more competitors in an area, it takes a shorter distance to physically browse the same number of suppliers.

[5] If search cost was to be independent of the number of competitors, then the expected price would rise with the number of competitors, $\frac{\partial E(p)}{\partial N}=\underbrace{\frac{\partial r}{\partial N}}_{>0}A+(r-c)\underbrace{\frac{\partial A}{\partial N}}_{>0}>0$. This is a result well established in the literature from the early work of Varian (1980) and Stahl (1989).

Furthermore, if the second derivative of the (expected) price with respect to the number of competitors is negative, $\frac{d^2E(p)}{dN^2} < 0$, then the price change is, in absolute value, lower for larger values of N. This generates the second hypothesis of interest.

Hypothesis 2. *The magnitude of the effect of exit or entry is higher the less intense competition is (the lower is N).*

2.2 Heterogeneous consumers

In line with Pennerstorfer et al. (2020) we augment the previous model of homogeneous consumers, to allow for consumers to be heterogeneous and have a proportion ϕ of the population with a willingness to pay v with probability and the remaining $1 - \phi$ with a willingness to pay 0. After expressing $\tilde{\mu} = \phi\mu$ it is possible to rewrite the mixed strategy symmetric equilibrium price distribution in (1) by replacing μ with $\tilde{\mu}$. Then, provided that $\frac{\partial E(p)}{\partial N} < 0$ and $\frac{\partial^2 E(p)}{\partial N \partial \tilde{\mu}} < 0$ and given that $\frac{\partial \tilde{\mu}}{\partial \mu} = \frac{\partial \tilde{\mu}}{\partial \phi} > 0$ we can link informedness and willingness to pay, producing our fourth hypothesis.

Hypothesis 3. *Exit (entry) has a larger positive (reduced negative) impact on the price in areas with (a) fewer informed consumers, and (b) more consumers with a limited valuation for the product (lower willingness to pay).*

We test part (a) of this hypothesis in our empirical exercise. Because we do not directly observe willingness to pay, we can use this relationship together with our empirical findings to infer about it.

Although the above theoretical framework does not explicitly incorporate income, we assume that both willingness to pay and the level of informedness are related to income. This is reflected in our empirical design, and how we interpret our empirical findings of income. For this reason, before we move on empirically testing Hypotheses 1-3, we organise our thoughts on how income affects how much consumers are willing to pay for motor fuel, and how much they search (how informed they are).

Motor fuel is a non-discretionary part of household expenditure. Such products (similar to rent or food) typically display significant distributional differences in that it constitutes a larger fraction of poorer households' expenditure.[6] Moreover, low-income households are the least likely to be able to switch to more expensive substitutes if the price of petroleum goes up. Regarding the relationship between income and search, Byrne & Martin (2021) provide a carefully constructed review of the relevant literature on this and concludes that most evidence points in the direction that low-income households engage less with the market. On the other hand, De los Santos (2018) finds that search duration decreases with income and is greater for retirement-age individuals. Nishida & Remer (2018) also find a positive relationship between search costs and income. Our results contradict this for two possible reasons: (1) in our specific case of Western Australia, the FuelWatch petrol retail price comparison website allows online price comparison - engaging with search this way imposes little extra search costs as long as the household has access to the internet, (2) people with higher income may drive more to work and therefore search becomes part of their commuting (without having to engage in search specifically).

3 Petrol retail markets and the data

The Australian petrol retail market is characterised by a small number of very large, and many fringe players. These can be divided into three distinct types. Refiner-wholesalers are vertically integrated retailers such as BP, Caltex, Mobil, and Viva Energy/Shell. This includes refiner–wholesaler controlled sites and independently operated but refiner–wholesaler branded sites. Large independent retail chains are independent retailers such as 7-Eleven, United, Puma Energy, and On The Run. Some supermarkets also sell fuel, such as Coles Express and Woolworths. At the country level, the combined retail market share (based on sales volume) of the large vertically integrated firms dropped significantly from over 80% in 2002, to under 40% in 2017. At the same time, the market share of supermarkets and independents increased substantially. Regarding the individual brands, Shell/Viva Energy (trading under Coles Express) and Woolworths (a supermarket) had a respective market share of 20-25% over our study period, followed by BP and Caltex, just under the 20% mark. The remaining sales volume is supplied by independent retailers.[7] Regarding the number of retail units in our sample, BP (260 stores) Caltex (254 stores), and Shell (216 stores including Coles) are the largest.

The main component of our data is daily prices at petroleum retail outlets in Western Australia, for the period 2001-2019, which was downloaded from FuelWatch.[8] FuelWatch is a price comparison service to motorists in Western Australia. At the time of introduction, the website was a response to policy concerns about large levels of price dispersion in the

[6] For the UK, Mattioli et al. (2018) finds that the poorest households often spend around 20% of their income on motor fuel, and frequently more than this.

[7] `https://www.accc.gov.au/system/files/Petrol-market-shares-report.pdf`

[8] `https://www.fuelwatch.wa.gov.au`.

country, implying that some consumers would have been paying largely over the odds. Byrne et al. (2018) offers a detailed description of the FuelWatch data, here we focus on the most important features that are relevant for this paper. Since its launch in 2001, the scope of FuelWatch was largely extended in 2003, and today it covers approximately 80% of regional and 100% metropolitan retail outlets in Western Australia. This includes information about the geographical location of the retail outlet (precise address), the brand of the operator, and prices for unleaded petrol (ULP), premium unleaded petrol (PULP), 95 RON (octane) petrol, diesel, branded diesel, and liquefied petroleum gas (LPG). Not all outlets sell the whole range of products. The way consumers can access FuelWatch has also significantly improved since its launch and has been used to plug into various smartphone apps going back to 2010. Through FuelWatch, consumers have free access to the next day's petrol prices at the petrol station of their choice, reducing switching costs for consumers who use the internet to search for the best petrol deals.

We had 15,638,524 observations of daily petrol station level data for all products at 1299 (1053) petrol stations.[9] Most of these are ULP (3,964,180) and diesel (3,811,105) prices. As we are not directly interested in the daily variation of prices, and also to eliminate issues from rogue missing observations, we averaged the price data at the weekly level and limited our focus to the two most popular products, ULP, and diesel.

We also collected weekly average wholesale prices for the Western Australian region. We acknowledge that not all retail units pay the same wholesale price. Vertically integrated companies have better distribution systems and lower costs than independents. However, this brand-level wholesale price information is not publicly available. Instead, in our estimates, we will control for the brand to account for this cost variation. The wholesale price data was only available from 01 January 2004 onwards, which further reduced our sample size, leaving us with 489,721 observations of weekly petrol station level prices. Figure B.9 in the Appendix shows the over-time variation of ULP and diesel prices. There is significant seasonality in the data. As our interest is in the immediate shocks as a result of exit/entry, we de-seasonalised (removed weekly and yearly seasonality) the price data.

For each petrol station, we acquired the longitude and latitude coordinates (using its address) and applied the Haversine formula[10] to identify which petrol stations are located within 1, 2, and 5 miles from each other. To give an example, Figure 1 shows the number of competitors for a selected independent petrol station had within a 1, 2, and 5-mile radius in our observed study period. Not all of these petrol stations were always available and competing in the entire sample period. Figure 2 illustrates this for the same independent petrol station. It shows that on 1 Jan 2004 it had 4 active rivals (Mobil, BP, Ampol, Puma). Then in early 2004 Ampol exited the market. In 2008 Puma also left the market, and in 2010 BP also left. Later in 2010 BP, and soon after that Ampol re-opened.

Two important features of the data need to be introduced here. First, we sometimes have observations of short spells of "market exit" (i.e. a rival has no price observation for a short period). The data does not reveal whether there is a temporary gap in reporting the data or a genuine temporary closure of one of the stations (for example for restoration or development work). We drop these periods from our sample. Second, looking only at the number of competitors masks information about the identity of the competitor (in this case the independent had 5 BP stations competing within a 5-mile radius, therefore the finding that the BP station within 1 mile from the independent station decided to exit may have to be interpreted in this context.) We deal with this by introducing variables such as the number of same brand competitors in the area, or a vertical chain dummy.

We collapse the petrol station level data by its address and postcode and link it to local area codes (Statistical Area Level 2 code).[11] There are 137 distinct SA2 areas in our sample, which include 195 distinct postcode areas. SA2s generally have a population range of 3,000 to 25,000 persons, with an average population of about 12,000 persons in our sample. SA2s in remote and regional areas generally have smaller populations than those in urban areas. Using the SA2 code, we then link each petrol station to local characteristics, using data from the 2016/2011/2006 Australian censuses, and the Personal Income in Australia report of the Australian Bureau of Statistics (ABS). We match this data with the corresponding petrol stations and the prices reported for each corresponding census data (prices before 2008 were matched to the 2006 census, prices between 2009 and 2013 were linked to the 2011 census, and prices from 2014

[9]The 1299 stations include 1053 distinct forecourts. For some of these, there were ownership changes in our sample period, which is why we have 1299 distinct petrol stations.

[10]The distance over the earth's surface. We could have used more sophisticated distance measures (e.g. the driving distance between two places) but our objective was not to precisely estimate the relationship between travelling distance and shopping behaviour, rather look at how a change in the number of petrol stations in proximity of a petrol station affects prices.

[11]`https://www.abs.gov.au/websitedbs/d3310114.nsf/home/australian+statistical+geography+standard+(asgs)`

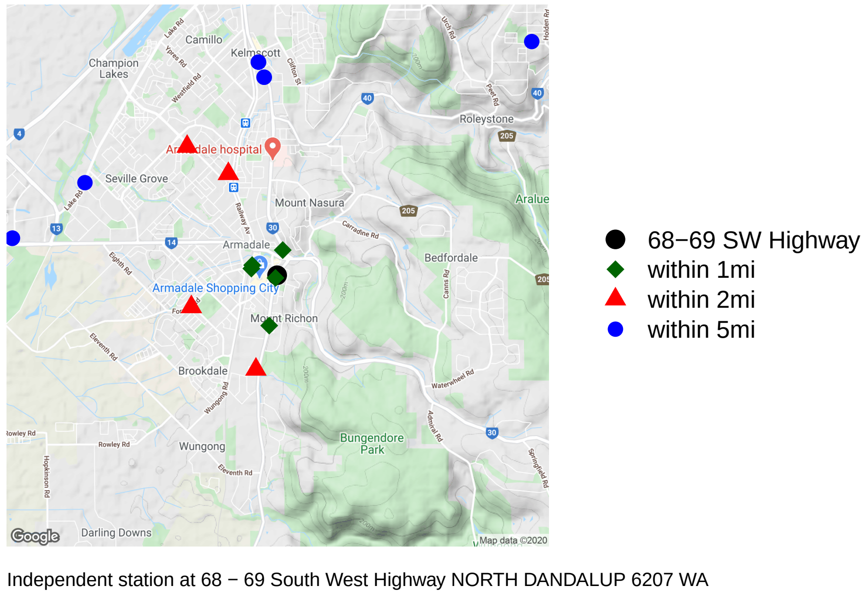

Figure 1: Example petrol station and surrounding competition

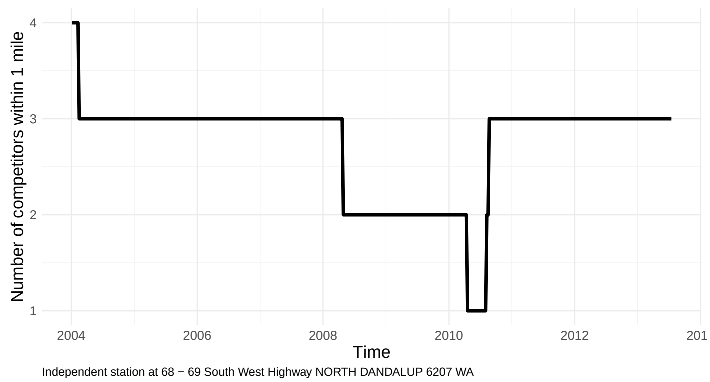

Figure 2: Number of competitors within a 1km radius for a petrol station in our sample

onwards were linked to the 2016 data.[12] The list of all our variables, and their summary statistics are given in Tables B.1 and B.2 in the Appendix.

Income: We record the annual taxable income at the postcode level, for the whole study period 2004-2019. Unlike some of the other area characteristics, this data is time-variant (annual). The data is collected from the taxation statistics of the Australian Taxation Office[13]. For individual income, the coverage is complete for 2004-2018. For business income data we have net business income for 2004-2018, and net rent for the same period. For the other business income variables, data is available between 2011-2018. For these, for the period 2004-2010, we assumed the same value as in 2011.

Search: We measure search through two main variables, the median commuting distance in an area, and the level of home internet penetration in the area. These two measures account for two dimensions of search: (1) people who drive more to work have a lower opportunity cost of search, as they already survey the prices as they drive past them; (2) people with home internet access are more likely to engage with the FuelWatch price comparison tool. Both can be

[12]The reason we did not use an exact matching is that demographic features are unlikely to change right at the time of taking the census. For example the local demographic characteristics in 2015) are likely to be better represented by the 2016 census rather than the 2011 one.

[13]`https://data.gov.au/data/dataset/taxation-statistics-postcode-data`

thought of as different dimensions of search costs. Internet access allows consumers to search online. In households without internet access, search has to be physical, which is associated with higher costs. Several previous works have looked at Internet use as a potential proxy for consumer informedness (Brown & Goolsbee 2002, Tang et al. 2010, De los Santos et al. 2012, Sengupta & Wiggins 2014). To also incorporate the cost of physical search, we measure commuting, which reduces the cost of search: the longer someone commutes to work, the more petrol stations they sample without incurring extra search costs. Previous empirical works that draw on information on commuting patterns to study heterogeneity in petrol prices include Cooper & Jones (2007), Houde (2012), Pennerstorfer et al. (2020).

ABS indices: The Australian Bureau of Statistics introduced several indices to measure the economic and social conditions in an area. The Index of Relative Socio-economic Disadvantage is a general socio-economic index. A low score indicates a relatively greater disadvantage in general. For example, an area could have a low score if there are many households with low income, or many people with no qualifications, or many people in low-skill occupations. The Index of Education (the level of qualification achieved or whether further education is being undertaken) and Occupation (classifies the workforce into the major groups and skill levels) reflect the educational and occupational level of communities. This index does not include any income variables. The Index of Economic Resources is a proxy for the financial aspects of relative socio-economic advantage and disadvantage (it summarises variables related to income and wealth). This index excludes education and occupation variables.

Other area characteristics: We have data on the age structure in each SA2 area (age, and % of people in various age brackets), the level of education, the average, and the mean commuting distance, and the means of commuting.

Year and quarter dummies: In an event study design, treatments in different calendar times are compared. Although we have removed annual and weekly seasonalities, the treatment may affect the retail price margin differently in different periods. For example, FuelWatch was designed as a price comparison tool in 2001, but consumers only gradually learned to use it over the years.

Brand: As we have a homogeneous product, there should not be much quality variation in the actual product, but there might be in the services linked to the product. The same good sold in two different stores could also be differentiated by the retail environment in which it is sold, with ‘high-quality’ stores charging more. Controlling for brands allows one to control for quality variation noise in our price variation data. Implicitly this assumes that heterogeneity in quality might exist across, but not within brands.

Note that we have much more diverse data on demand-side factors than supply-side ones. However, even in the absence of firm-level data, we can allow for supply-side heterogeneity by controlling for area-level supply-related factors, such as the average (by business) business expense in an area, the average business tax paid, the average business income, and the average rent paid by businesses.

4 Descriptive analysis and study design

4.1 Descriptive analysis

As shown in previous works, there can be substantial dispersion in the price of petrol (Pennerstorfer et al. 2020, Byrne & de Roos 2017). This is no different in our sample. To get a better understanding of the source of dispersion in our data, this section presents some descriptive information on the retail margin. First, Table 1 shows how the retail margin varies with the level of competition. The table confirms conventional IO theoretical and empirical evidence that higher market concentration is associated with higher margins.

Table 1: Margin by competition

number of rivals within 1mile	**0**	**1**	**2**	**3**	**4**	**5**	**6**	**7**	**8**	**9**	**10**
ulp margin	1.134	1.121	1.112	1.124	1.118	1.118	1.115	1.105	1.094	1.085	1.092
diesel margin	1.134	1.129	1.123	1.132	1.123	1.125	1.121	1.115	1.126	1.113	1.102

Table 2 provides more insight into how the margins differ around our main variables of interest: competition, income, and our two measures of search costs, internet (based on the % of households with home internet access) and commuting distance (based on the distance commuted to work). In Table 2 we compare the margins for areas with low and high levels of competition, income, commuting distance, and internet access (to define low and how we split our sample around the median values of these four variables). The numbers confirm, that lower competition is associated with higher margins. Our theoretical introduction posited that some of the price dispersion may be due to the heterogeneity in the level of engagement with the market. In our measurement of search, low commuting distance paired with low internet access in an area is assumed to be associated with the lowest engagement (highest search costs), and vice versa. Table 2 reveals that areas with the highest proportion of informed consumers face the lowest margins. Regarding income, there is a mixed picture, in low competition areas low-income areas are associated with lower margins, but the same is not true for high competition areas. Table B.3 in the Appendix shows that averaging across the total sample, low-income areas experience slightly higher margins in general. Margins are also higher in low education areas, areas with less home internet penetration, and places with a higher proportion of people over 65.

Table 2: Mean margin by levels of competition, income, and search

		ULP			
		low internet		high internet	
		low commute	high commute	low commute	high commute
low competition	low income	1.188 (0.091)	1.124 (0.057)	1.154 (0.065)	1.122 (0.061)
	high income	1.208 (0.096)	1.146 (0.084)	1.197 (0.097)	1.094 (0.047)
high competition	low income	1.13 (0.067)	1.077 (0.027)	1.1 (0.039)	1.09 (0.033)
	high income	1.084 (0.037)	1.078 (0.03)	1.086 (0.031)	1.084 (0.032)
		Diesel			
low competition	low income	1.17 (0.075)	1.123 (0.043)	1.142 (0.056)	1.125 (0.056)
	high income	1.184 (0.076)	1.136 (0.055)	1.169 (0.078)	1.109 (0.044)
high competition	low income	1.135 (0.051)	1.102 (0.024)	1.111 (0.037)	1.109 (0.039)
	high income	1.11 (0.031)	1.105 (0.029)	1.12 (0.035)	1.112 (0.039)

Moreover, as Table B.6 in the Appendix shows, there is also significant variation in the price margin across the brands. The small independent stores operate with the highest margins, but the large vertically integrated companies (BP, Caltex, Shell/Coles) are also in the top third. The bottom half of the distribution (lowest margins) constitutes mainly independent chains. This would suggest that cost-efficiency is likely to be dominated by other factors when it comes to setting the margin, as vertically integrated companies are likely to have lower retail costs, but still choose to have a high margin. Some of these differences may be explained by local cost conditions. Stores in urban areas may face higher rental prices and labour costs, which would lead them to charge more, though they might also face more intense local competition, leading them to charge less. Independent ‘corner’ shops are unable to exploit economies of scale in wholesale purchasing or other costs and so charge higher prices than large, national retailers for the same product. If poor households are concentrated in areas with high retail costs, then this may explain any finding that the poor pay more.

It is also possible that firms behave strategically when choosing whether to open stations in rich or poor areas. Table B.8 in the Appendix shows that this does not seem to be the case in our data. Looking at the largest brands (BP, Caltex, Shell), we can see that low-income areas often have more competitors, but at the same time also higher margins. To further confirm this, Table 3 shows that when the market is defined as a 1, or a 2-mile radius, the number of rivals is similar in low and in high-income areas (there is a difference when one looks at the significantly wider 5-mile radius geographical market). Moreover, there seems to be a difference in consumer informedness, with high-income areas displaying signs of more informed consumers.

Table 3: Main data features by income groups

	N within 1mi	N within 1mi per 10000 people	N within 2mi	N within 2mi per 10000 people	N within 5mi	N within 5mi per 10000 people	internet	commute
high income	2.831 (1.751)	3.656 (2.864)	5.539 (3.771)	5.973 (3.574)	16.288 (17.811)	14.651 (12.85)	0.239 (0.06)	6.457 (4.152)
low income	2.911 (1.748)	2.761 (2.138)	6.817 (4.425)	5.863 (4.006)	29.351 (22.43)	22.786 (16.253)	0.264 (0.062)	7.666 (4.664)

Although these descriptive tables are useful for understanding the data, to test our hypotheses, we need an approach that brings together all possible effects into the same model. This is what we set up with our study design.

4.2 Study design

We estimate the causal impact of the exit and entry of petrol retail forecourts on the retail margin and investigate the heterogeneity of these estimates across different area (consumer) characteristics. Because we do not observe the same markets with and without exit/entry at the same time, we rely on observational data, which requires us to design a carefully selected control group that can be used as a stand-in for the outcome that would have happened in the absence of exit/entry.

We employ an event study design to line up all relevant events (exits and entries). The use of event study design in quasi-experiments is increasingly common, Schmidheiny & Siegloch (2019) point out that around 5% of the papers published in QJE, AER, and JPE in the most recent years used an event study design. There are also numerous recent methodological contributions, such as Freyaldenhoven et al. (2019), Sun & Abraham (2020), Roth (2018), or Borusyak & Jaravel (2017). The research design implies that for our study, absolute time is normalised such that the observation period is measured concerning treatment time (observations are lined up around the treatments).

We created an event window (15 weeks before and 15 weeks after the treatment), which we applied to every petrol station that experienced exit/entry. We defined exit as the reduction of the number of rivals a petrol station has within a 1-mile radius (in the Appendix we provide our main results for exit/entry within 2 and 5-mile radii). We removed from our sample all instances of exit and entry, where there was another exit or entry within the same area, in our pre, and post-treatment periods (26 weeks before and after the treatment). This ensured that there was no confounding effect from another change in local market structure well before and after the treatment. The reduced sample included 392 instances of exit and 354 instances of local market entry. Figure 3 shows the annual distribution of these exit and entry events. There is a reduction in the number of exits, partly because with fewer petrol stations in the market, there are fewer potential stations to exit. Entries happen at a roughly even rate over time. The figure also displays the quarterly breakdown of exit and entry. This shows an increased number of exits and entry in the second and third quarters of the year, which is likely to do with the end of the tax year (end of June).

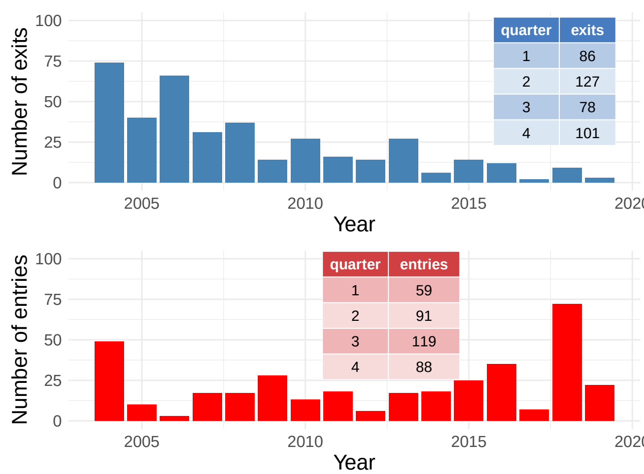

Figure 3: Number of entries and exits in our working sample

The frequency of entry and exit also varies across the different areas. Table 4 shows the ratio of exits and entries to the total number of petrol stations, broken down by our four main variables of interest (competition, income, and search costs as measured by internet access and commuting distance). With low competition, high-income areas witness proportionately more exits. High-population areas are more likely to see more changes in market structure. In general, it appears that areas with more competitors also see more shifts in market structure.

Table 4: The the ratio of exits and entries to the total number of petrol stations by income, competition and population size

		low competition		high competition	
		low population	high population	low population	high population
Exits	low income	0.295	0.344	0.398	0.75
	high income	0.404	0.508	0.451	0.312
Entries	low income	0.299	0.296	0.341	0.38
	high income	0.34	0.516	0.28	0.359

In our event-study design, we align these instances of local market exit and entry for each of our ULP and diesel data. This creates unbalanced panel datasets of treated and units for the years 2004-2019 with varying dates of treatment application. Each instance of treatment (exit and entry) is therefore defined as petrol stations that had another station exit or enter the market within a 1-mile radius. Our control group draws from all other petrol stations that did not experience a change in the number of rivals ± 26 weeks from the time of the treatment. Studies with a similar research design, that are only interested in the average treatment effect often average over all these potential control units. In our case, we are interested in the heterogeneous treatment effect. Averaging the control units would mean averaging their characteristics, making them ill-suited for our purpose. Instead, we decided to take the most similar petrol stations (based on the observable features in the period before -15 weeks from the exit/entry). For this we employed nearest neighbour matching, using the propensity score difference to specify distances from the treatment petrol station, and selecting the petrol station with the lowest distance. In a set of experiments, we looked at the nearest 1, 5, 10 neighbour(s) (each treatment petrol station is matched with 1, 5, 10 control petrol stations) but in our main discussion, we focus on the nearest 5 neighbours. We offer a sensitivity analysis of this choice in the Appendix. This gives us, for each instance of exit and entry, a set of 6 petrol stations (1 treatment and 5 control). Table B.4 in the Appendix compares the average and the standard deviation of the treatment and control groups to demonstrate similarity on observables.

In observational studies, a cursory look at the raw data can often provide a useful indication of whether the testable hypotheses hold. Put differently, caution is warranted if the answer to the main research question is not apparent from simple descriptive figures. Figure 4 shows how the average retail margin varies for the treatment and control groups for ULP for the lowest and highest income terciles. The vertical lines represent the time of exit. There seems to be an increase in the margin, but only in the low-income areas.

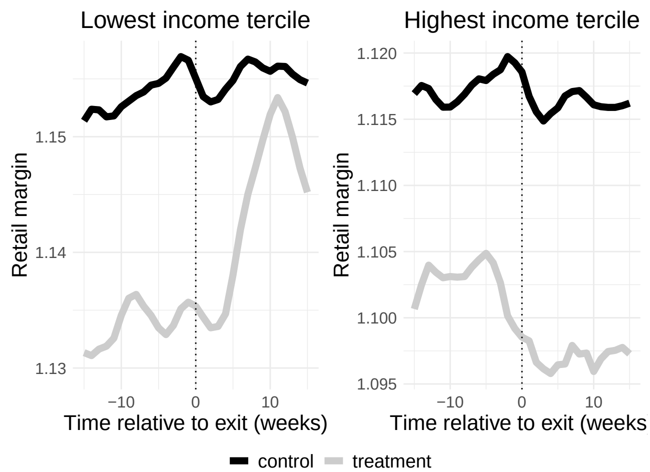

Figure 4: Retail price margin before and after exit for ULP at different income levels

Figure 5 shows that following an entry, the treatment group experienced a clear drop in margins for both low and high-income areas. Both of these figures confirm conventional industrial organisation theory and previous empirical findings, i.e. more competition leads to lower prices, although there seems to be a difference in the level and the change in the level of margins between low and high-income areas. In the following section, we formally test this difference.

5 Econometric method and main results

5.1 Estimating heterogeneous effects

We need to estimate the treatment effect of exit and entry on the retail margin. The conceptual problem is similar to that formulated in Rubin (1974). Denote a vector of covariates for petrol station i by X_i. We let the treatment (exit and entry from and to the market) indicator W_i take on the values 0 (the control group, i.e. no exit/entry) and 1 (the treatment group, i.e. exit/entry). For petrol station i, $i \in 1, ..., N$, let Y_i denote the observed outcome and the outcome of interest (the retail margin) in the case of receiving the treatment as $Y_i(1)$ and when not receiving the treatment as $Y_i(0)$. The causal effect of exit and entry for petrol station i is therefore $Y_i(1) - Y_i(0)$. The Conditional Average Treatment Effect (CATE) is given by:

$$\tau(x) = E[Y_i(1) - Y_i(0)|X_i = x] \tag{5}$$

The problem of causal inference is that we do not observe both $Y_i(1)$ and $Y_i(0)$ at the same time. Instead, we estimate CATE by a difference in means $y_t - y_c$, where y_t and y_c are the means of the outcome variable for the treated and control groups, respectively. For identification, we assume unconfoundedness, $\{Y_i(1), Y_i(0)\} \perp W_i|X_i$, i.e. the markets that experience exit/entry are selected randomly conditional on the covariates.

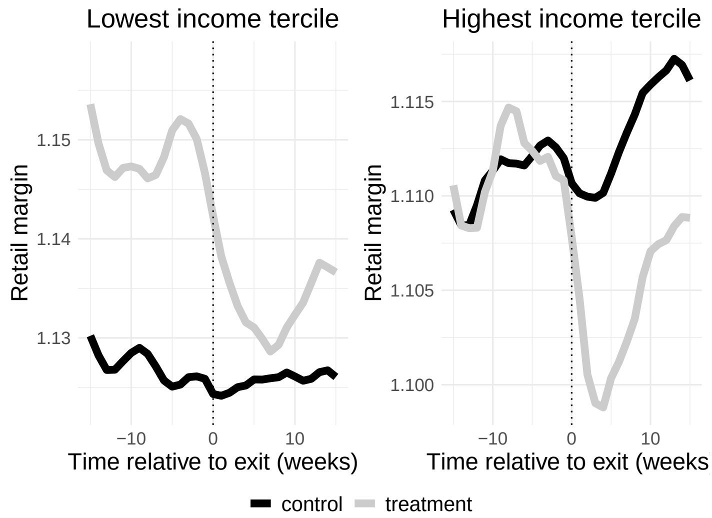

Figure 5: Retail price margin before and after entry for ULP at different income levels

Our objective is to conduct estimation and inference on the function $\tau_i(x)$ to gain insight into the heterogeneity of the treatment response, across our observable local area characteristics. One way to do this would be through introducing interaction terms in the estimation of $\tau_i(x)$. Alternatively, we could estimate $\tau_i(x)$ for different sub-samples of the data. The problem with these solutions is that they run the risk of using a misspecified model, and even if the correct model was estimated, it can quickly run into issues of dimensionality.[14]

To move beyond the above solutions, numerous recent methods have been proposed for estimating heterogeneous treatment effects: Imai and Ratkovic (2013), Signorovitch (2007), Tian et al. (2014), Weisberg and Pontes (2015); Beygelzimer and Langford (2009), Dudík, Langford, and Li (2011). In the econometrics literature, Bhattacharya and Dupas (2012), Dehejia (2005), Hirano and Porter (2009), and Manski (2004) estimate parametric or semiparametric models for optimal policies, relying on regularization for covariate selection in the case of Bhattacharya and Dupas (2012).

In this paper we use generalised random forests (causal forests) as proposed by Wager & Athey (2018) and Athey et al. (2019). The method fits our problem for multiple reasons. As non-parametric a tree-based method, it does not require us to specify a (potentially complex) linear relationship between our covariates and the treatment effect. It also allows the efficient handling of large covariate spaces. In our case the number of possible features (accounting for all interactions) is much larger than the sample size, therefore methods such as OLS cannot be considered without the research filtering which features to use first. Moreover, Athey et al. (2019) showed that the estimates achieve asymptotic normality and as such, it is suitable for hypothesis testing on the treatment effects. Finally, independent variables are time-invariant in our event-study setup (they change every 1 or 5 years). Estimating a 2-way fixed effects linear model would mean simply losing our covariates as they would be subsumed by the fixed effects dummies.

Below we provide a brief introduction to tree-based methods and causal forests. This should be sufficient for those who are unfamiliar with these methods to understand its intuition, but for details, we refer the reader to the literature cited above.

Regression trees are a non-parametric machine learning approach, which are frequently used for prediction problems in data science.[15] Assume we have k covariates and N observations, and we want to partition the covariate space $\mathcal{X}$ into M mutually exclusive regions $R_1, ..., R_M$, where the outcome for an individual with covariate vector X in region R_m is estimated as the mean of the outcomes for training observations in leaf R_m. Denote the subset corresponding to R_m as $\mathbb{X}_m$. Let X_j be a splitting variable and s be a split point. Define $R_1(j,s) = \mathbb{X}_1 = \{X \mid X_j \leq s\}$ and $R_2(j,s) = \mathbb{X}_2 = \{X \mid X_j > s\}$. The algorithm selects the pair (j,s) that solves:

$$\min_{j,s}\left[\sum_{X_i \in \mathbb{X}_1}(Y_i - \overline{Y}_1(j,s))^2 + \sum_{X_i \in \mathbb{X}_2}(Y_i - \overline{Y}_2(j,s))^2\right] \tag{6}$$

[14]For k potential sources of heterogeneity, this would mean adding $2^k - k - 1$ interaction terms to our model.
[15]For details see: Breiman et al. (1984)

where $\overline{Y}_1(j,s)$, and $\overline{Y}_2(j,s)$ are the mean outcomes in $R_1(j,s)$ and $R_2(j,s)$. Eq.6 splits the data into two regions, then the process is repeated on each of the two resulting regions. Regression forests are ensemble methods, whereby the forest predictions are constructed as the average of the tree-based predictors. (Eq.6) can also be thought of as the 'growing' or 'splitting' part of constructing regression trees.

Causal trees build on the same concept, but for each node, instead of minimising the mean squared error (MSE) for the difference between the average outcomes for each node, it minimises the MSE for the difference in the estimated treatment effects.

Athey et al. (2019) proposes using an honest approach to estimating these causal trees, i.e. they grow the tree on a sample of the data, and they estimate it using a different sample. In the context of causal trees, the idea is that the leaves are small enough that the (Y_i, W_i) pairs for each leaf had come from a randomised experiment. In this case, the treatment effect in the small space of each leaf with the corresponding set $\mathbb{X}_m$ is simply:

$$\hat{\tau}_{\mathbb{X}_m} = \frac{1}{|\{i : W_i = 1, X_i \in \mathbb{X}_m\}|} \sum_{\{i:W_i=1,X_i\in\mathbb{X}_m\}} Y_i - \frac{1}{|\{i : W_i = 0, X_i \in \mathbb{X}_m\}|} \sum_{\{i:W_i=0,X_i\in\mathbb{X}_m\}} Y_i \tag{7}$$

Finally, to construct a causal forest, we draw repeated bootstrap samples (b) of size N from the training data to recursively estimate a number of causal trees. The prediction for an individual with a vector of covariates X is then $\hat{\tau}_i = \frac{1}{B}\sum_{b=1}^{B}\hat{\tau}_b$ where $\hat{\tau}_b$ is the estimate produced by tree b. Athey et al. (2019) show that the estimated treatment effect is asymptotically normal.

Causal forests are useful for finding heterogeneity in the treatment effect in a cross-section setup. Our event study design, however, means that we have longitudinal data, so does this mean that we lose important information by our choice of method? The variables in X_i are area petrol station and area characteristics and can be considered constant within the 30-week event window of our analysis. The outcome variable Y_i on the other hand is time-variant. Therefore choosing causal forests as our method does mean that we forego the possibility of estimating time-dependent treatment effects (for example in the sense of traditional event study designs). We believe this trade-off is justified as we are primarily interested in the heterogeneity in the treatment effect, rather than its dynamics.

For our causal forest therefore the outcome variable of interest Y_i is the change in the retail margin for petrol station i before and after the exit/entry:

$$Y_i = \frac{1}{|\{T_1\}|}\sum_{t\in\{T_1\}} margin_{it} - \frac{1}{|\{T_0\}|}\sum_{t\in\{T_0\}} margin_{it} \tag{8}$$

T_0 and T_1 represent the pre- and post-treatment periods, respectively.

In X_i we include the features listed in Table B.1, which reveals some overlap. For example, we have many different ways of measuring education, or wealth, and we have no a priori knowledge, which one of these is important in driving the treatment effect heterogeneity. Athey & Wager (2019) proposes removing the least important features from the estimation of causal trees to improve estimates. This is a feasible option, but we are specifically interested in the effect of some variables on the treatment effect, and this solution may eliminate some of our variables of interest. Instead, we add an extra layer to causal forests with a bagging ensemble learning method. The idea is, to randomly draw several features, add our features of interest, and re-estimate the forest in each draw, on this reduced sample of features. This way the estimated ensemble treatment effects are:

$$\tau^G = \frac{1}{G}\sum_{g=1}^{G}\hat{\tau}_g \tag{9}$$

where G is the number of causal forests we run to get our ensemble individual treatment effects. The standard errors are derived from the bootstrapped standard errors of the individual causal forests and the squared deviation of the treatment effects:

$$\sigma_{\hat{\tau}^G} = \sqrt{\frac{\sum_{g=1}^{G}[\hat{\sigma}_g^2(\hat{\tau}_g - \bar{\tau})^2]}{G}} \tag{10}$$

We argue that this ensemble method is more fitting in cases where there is a relatively small sample size, and a large number of parameters, and we have specific (theory-driven) interest in a selected set of these features. In Section A of the Appendix, we provide details and simulations to justify our approach.

5.2 Causal forest results

Table 5 shows the conditional average treatment effects (CATE) and the conditional average treatment effects on the treated (CATT). The average effects are as anticipated in our descriptive part: exit, on average triggered a small (statistically not significant) increase in the margin, entry, on average lead to a larger drop in prices. Our interpretation of the asymmetry between exit and entry is to do with the level of market concentration in markets where we observed exit and markets where we were sampling instances of entry. Table C.1 in the Appendix shows that in our estimation sample entry was more likely to happen in more concentrated markets. In these markets, the effect of a change in market concentration on the retail margin is more pronounced.

Table 5: Conditional average treatment effects

	Exit		Entry	
	CATE	**CATT**	**CATE**	**CATT**
ULP	0.08 (0.065)	0.079 (0.058)	-0.329 (0.081)	-0.335 (0.072)
Diesel	0.082 (0.064)	0.086 (0.061)	-0.067 (0.077)	-0.059 (0.073)

Bootstrapped standard errors in parentheses.

However, the standard deviation of these average treatment effect estimates suggests that there can be substantial variation in the individual treatment effects across the petrol stations in our sample. To start exploring this heterogeneity, Figure 6 offers visual verification of the relationship between the income and the treatment effect for ULP and diesel. The red horizontal line shows zero treatment effect, and the black vertical line indicates the median value of the feature on the horizontal axis (income). The figure reveals that in areas below the median level of income, the treatment effect is very dominantly positive (and large). Above the median income, $\hat{\tau}_i$ is closer to and around zero. The income-related heterogeneity seems more pronounced for ULP than for diesel. In Section C of the Appendix, we provide more figures to show the relationship between several of our variables and the treatment effect.

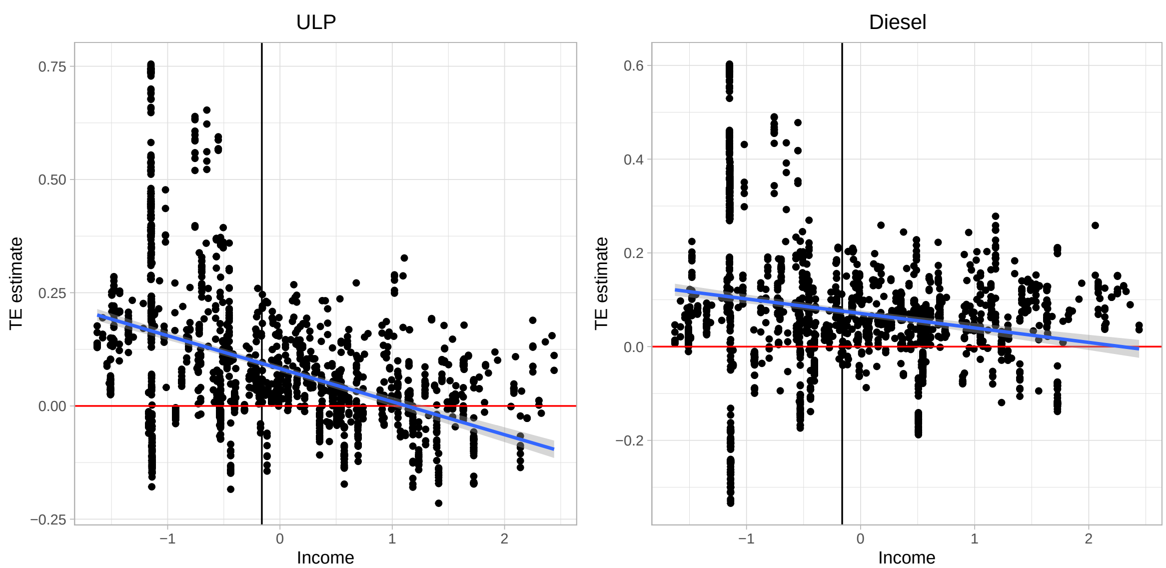

Figure 6: Treatment effects by income

To test our three hypotheses, Table 6 presents a more detailed breakdown of the treatment effects related to exit in ULP, breaking it down to our main variables of interest. We defined low and high for these four variables, by taking the values corresponding to their 10th and 90th percentiles respectively. We then used our estimated causal forests to predict the treatment effect, assuming mean values for all other covariates.

Several stylised findings can be deduced from this exercise. Most importantly for our investigation, with exit, lower-income households see a larger price increase. This difference is more pronounced in markets where competition is lower. The competition rows indicate that prices increase more in markets that were less competitive before the exit. This suggests that increasing market concentration increases price dispersion (the extent to which businesses choose to price discriminate) with low-income areas seeing a larger price increase. Moving on to our measures of the informedness of consumers, in general, estimates in the bottom right corner (which imply more informed consumers) are lower than estimates in the upper left corner for both low and high competition levels. It appears that commuting more, i.e. having a higher chance of physically browsing prices, reduces the price inflating impact of exit more than having more households with access to home internet. Areas, where people commute longer to work, experience lower price increases as a result of exit, as though more commuting was proportionate with search intensity. This is consistent with several previous works such as Pennerstorfer et al. (2020).

Table 6: Predicted treatment effects of **exit** in **ULP** by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.271	0.243	0.272	0.237
		(0.118)	(0.109)	(0.126)	(0.118)
	high income	0.125	0.097	0.116	0.079
		(0.086)	(0.076)	(0.09)	(0.082)
high competition	**low income**	0.127	0.113	0.13	0.109
		(0.074)	(0.065)	(0.08)	(0.07)
	high income	0.035	0.02	0.03	0.009
		(0.056)	(0.049)	(0.059)	(0.053)

Predicted treatment effects, "low" implies fixing the given variable at its 10th, and "high" refers the 90th percentile.
Bootstrapped standard errors in parentheses.

Entry on the other hand results in a fall in prices, which is more pronounced in areas with less competition. This is intuitive, on the margin, areas with low levels of competition can gain more from a new rival. We get a more peculiar result regarding the heterogeneity due to differences in income levels. Areas with different income levels experience similar price drops as a result of falling local market concentration (marginally larger in high-income areas). Both of our measures of consumer informedness suggest that more informed consumers are associated with a larger fall in the price margin following entry.

Tables C.4 and C.5 in the Appendix report the results for diesel. The tables tell a similar story (somewhat smaller magnitude) for the diesel margin following exit. The effect of entry is negative, although much smaller (in magnitude and significance) than in ULP. The role of income and search is similar to ULP - areas with more high-income households (or more informed households) experience a smaller diesel margin increase with exit. With entry, there does not seem to be a pronounced difference between the price drop experienced by low and high-income households.

Table 7: Predicted treatment effects of **entry** in **ULP** by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	-0.264	-0.251	-0.366	-0.346
		(0.045)	(0.039)	(0.053)	(0.048)
	high income	-0.306	-0.294	-0.381	-0.364
		(0.041)	(0.037)	(0.044)	(0.04)
high competition	**low income**	-0.223	-0.212	-0.329	-0.31
		(0.036)	(0.031)	(0.047)	(0.042)
	high income	-0.268	-0.258	-0.346	-0.329
		(0.032)	(0.029)	(0.039)	(0.035)

Predicted treatment effects, "low" implies fixing the given variable at its 10th, and "high" refers the 90th percentile.
Standard deviation in parentheses.

Altogether, these results suggest two main effects that hold for each product and both exit and entry. The number of rivals in the market is important, concentrated markets witness a larger rise/fall in the price margin from exit/entry. The results on income are also consistent, low-income households always experience a larger increase in the price margin with exit. The impact of the informedness of consumers remains similar across the two products. More informed consumer in an area is associated with a lower price increase with exit and a larger price drop with entry (with the exception of entry in diesel).

6 Identification assumptions and robustness

6.1 Unconfoundedness assumption

Up to this point, we have assumed that selection to treatment was random (unconfoundedness assumption), conditional on our observable variables. Non-random selection means that unless all relevant variables are observed, our estimates will be biased. A frequent violation of the random assignment assumption is when unobserved factors are correlated with the treatment and the outcome variable in question, leading to biased estimates (omitted variable bias). Conventionally, researchers try to remedy this problem by employing fixed effects, or instrumental variables in their models. The problem with this approach is that it relies on strong assumptions that may not hold, and it is hugely limited by dimensionality issues in a conventional linear regression setup. For example, in our study, some convoluted (non-linear) interactions between the independent variables likely affect the treatment, and not controlling for this would lead to biased estimates. But the use of a linear model constrains researchers in how many of these interactions they can include in their models. Our choice of method handles this problem and allows a much richer set of observable factors to control for. Although it is never possible to observe and account for all relevant factors, under our model the conditional independence assumption relies on a much wider range of attributes than would be possible in linear models. We can include a large number of observed variables and their interactions with the way the treatment affects the outcome, reducing the risk of omitted variable bias.

It is also possible that not the treatment, but an event before the treatment consistently confounds our estimates. To test this, we look at pre-treatment parallel trends by focusing on the 6 months preceding the treatment. If pre-treatment the parallel trend assumption is not violated, we would expect to see zero treatment effect. For this exercise we assumed a placebo treatment to happen halfway through our test period (3-months before the real treatment). Table C.8 shows our results for the impact of exit on the ULP margin and finds that the effect is not significantly different from zero in any of the tested instances.

6.2 Exogeneous treatment

There is also a possibility, that market structure is endogenous (reverse causality), although numerous previous studies treated petroleum market exit and entry as exogenous in the retail markup equation (Simpson & Taylor 2008, Hastings 2004, Taylor et al. 2010). Reverse causality would mean that not only exit/entry would affect the retail margin, but the retail margin would drive exit/entry, something that sounds intuitively plausible. But in our study design, this should only concern us if the margin in the vicinity of exit/entry is what drives the decision to exit/enter. Given the cost of exit/entry, it is highly unlikely that these decisions are made on the whim of short-term fluctuations in the margin.

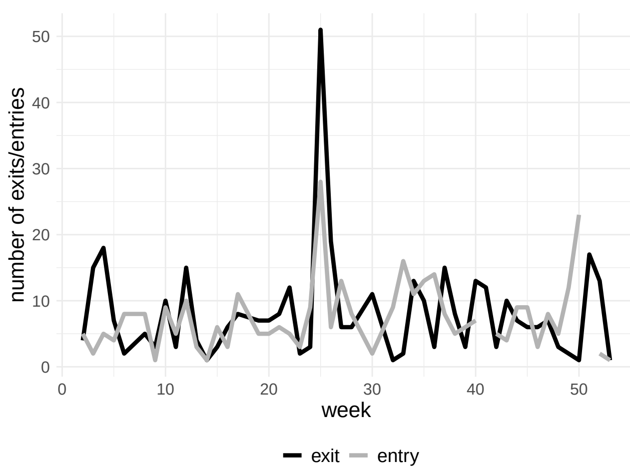

Figure 7: The number of exits and entries by calendar week

This is supported by Figure 7, which shows a large spike in the number of exits and entries at week 25, which is the end of the tax year in Australia. For administrative purposes, it makes sense for businesses to close down at the end of the tax year, or open up right at the beginning of the new tax year. This implies that at least a large number of exit and entry instances in our data were not hastily made in response to a short-term change in the retail margin and therefore can be considered as exogenous in our study design.

Table 8: Comparing the samples around the end of tax year (ETY)

		income	internet	commuting	comp1mi	comp2mi	comp5mi
exit	**month before ETY**	46164.32	0.229	6.286	2.751	6.393	22.028
		(9219.653)	(0.062)	(4.399)	(2.135)	(4.610)	(23.066)
	rest of the sample	47455.648	0.226	5.997	2.914	6.538	22.653
		(10087.512)	(0.072)	(4.353)	(2.312)	(4.916)	(22.292)
entry	**month after ETY**	47101.253	0.225	6.059	3.158	7.103	23.544
		(9842.299)	(0.083)	(4.198)	(2.238)	(4.971)	(22.805)
	rest of the sample	47299.158	0.227	6.033	2.873	6.477	22.500
		(9997.499)	(0.070)	(4.371)	(2.293)	(4.867)	(22.365)

To explore the strong possibility that these exits and entries around the end/start of the tax year are exogenous, we re-estimated our causal forests for the subgroups of exits happening within 4 weeks before the end of the tax year, and subgroups of entries happening within 4 weeks after the start of the tax year. These respective sub-samples contain 73 instances of market exit and 35 instances of entry. Firstly, Table 8 compares the mean and standard deviation of our main variables of interest for these sub-samples with the rest of the sample. There appears to be no systematic difference between these exits/entries and the rest of the sample.

Table 9 shows the conditional average treatment effects for these sub-samples. We only replicate estimates of the conditional average treatment effects and not the estimates on the heterogeneity in the treatment effects, because here we have a limited sample size with more limited variation in the features that we expect to drive this heterogeneity (competition, search, income). The CATEs reported in Table 9 are similar to those estimated for the total sample (Table 5). Standard errors are higher, due to the limited sample size. If one accepts the claim that these exits/entries around the end of the tax year are exogenous, this finding suggests that our main results are not biased by potential endogeneity (reverse causality).

Table 9: CATEs for the samples around the end of tax year for ULP

exit		**entry**	
CATE	CATT	CATE	CATT
0.056	0.082	-0.247	-0.368
(0.148)	(0.134)	(0.252)	(0.217)

Bootstrapped standard errors in parentheses.

To provide further reassurance that reverse causality is not affecting our findings, we provide a set of experiments, where we narrow the pre-event side of our study window. The idea is to move our study window closer to the treatment date, to ensure that changes in the margin were not affecting the decision to exit/enter (for example a change in the margin 5 weeks before the exit is unlikely to be the reason for exit). Tables C.9 and C.10 in the Appendix show that our qualitative results remain, although the magnitude of the results changes.

6.3 Robustness checks

In the results highlighted above, we looked at how the exit or entry of a rival within a 1-mile radius impacts the ULP and diesel margins of a petrol station. In our experiments, we also looked at what happens to the same margins if a rival from a 2-mile and a 5-miles radius exits or enters the market. The results for ULP exit (within 2 miles and within 5 miles) are presented in Tables C.11, and C.12, respectively. Both tables follow the same logic as 6 above. We use the estimated causal forests to predict the treatment effect for various levels of competition, income, internet, and commuting. High levels refer to the value of the respective variable at the 90th percentile, and low refers to the value at the 10th percentile. The tables show that the treatment effect falls as we are looking at the impact of a petrol station exiting/entering at a further distance, with the largest effect size from an exit/entry within 1-mile (Table 6), lower average treatment effect if the exit/entry happened within 2-miles away, and even lower if it happened 5-miles away. This is expected. The exit/entry of a close rival is likely to have a larger absolute impact on the retail margin.

As with other tree-based methods, causal forests allow the clustering of estimands (see Athey & Wager 2019). We use this as a robustness check to cluster the petrol stations that experience the same exit/entry as one of their competitors. Results with causal forests clustered by postcode are reported in TablesC.13 and C.14 in the Appendix.

Finally, we also looked at how sensitive our results are to choosing a different nearest neighbour matching to select our control group petrol stations. Tables C.15, C.16 in the Appendix show the results for choosing the 2 and the 10 nearest neighbours. Our story remains qualitatively unchanged.

7 Discussion of the results

Firstly, we presented evidence supportive of exit leading to a small increase, and entry triggers a larger drop in the price margin. We argue that the asymmetry is because our sample entry tends to happen in more concentrated markets. Although the effect of exit is not significant on average, by looking at treatment effect heterogeneity, we identify the cases where it leads to a significant increase in the margin. First of all, the margin increasing effect of exit, and the margin reducing the effect of entry are larger in absolute value in less competitive markets, meaning that we cannot reject Hypotheses 1 and 2.

We also offered detailed evidence that low-income households are punished with a larger increase in the retail margin of petroleum products when market concentration increases. At the same time, we do not find that the same low-income households enjoy a larger drop in margins when competition intensifies (in fact high-income households seem to enjoy a somewhat larger drop in the margin). Starting from the theoretical and empirical results of search literature, which suggests that heterogeneity in the level of engagement with the market can lead to price dispersion even in homogeneous goods, our results imply that low-income households may be engaging less, and as a result, changing market concentration can have distributional effects.

Looking at the role of search (and confirming part (a) of Hypothesis 3) we found that in areas with a larger proportion of informed consumers there was a lower increase in prices as a result of exit and a higher fall in margins as a result of entry. Interestingly though, there was still a difference between low and high-income households even after accounting for these measures of search and informedness. In our models, we control for a large number of observable features and income remains an important factor in how much margins increase/fall as a result of increasing/falling market concentration. For example, age and education are often-cited sources of search heterogeneity. But Tables C.6 and C.7 show that even after fixing the level of age and education, the difference between low-income and high-income households remains. This would suggest that unobservable factors may play an important role in how low and high-income households engage with the market. Byrne & Martin (2021) for example argues that differences on a cognitive level, differences in biases, and in how people process information may also be behind low-income people engaging less with the market.

It is also likely that low-income households have a higher willingness to pay. This statement may seem counter-intuitive at first glance, but it makes sense if one considers that higher motor fuel prices eventually encourage the average consumer to cut back on driving or switch to more fuel-efficient vehicles, however in the short-run low-income households may have few options but to continue buying motor fuel and cut back on other expenditures (or get further into debt). Our finding that low-income households experience a larger retail margin increase after exit even in high search areas, suggests that their higher willingness to pay may also play a role.

Of our search variables, commuting seems to play a more important role in improving consumer informedness than internet access. Several previous works argue that the Internet and price comparison websites do not contribute to better-informed consumers (Ellison & Fisher Ellison 2005, Ellison & Ellison 2009), and our findings may be interpreted in support of these arguments. But it is also important to add that our measure of the Internet is through the level of local penetration of home Internet, which therefore does not capture access to mobile Internet, which has likely dominated Internet use in the second half of our sampled years.

There seems to be a difference between the ULP and the diesel results, which is not unexpected. Diesel demand is likely to be dominated by commercial users (heavy goods vehicles, and large goods vehicles). This would explain that there is less of an impact of changing local market concentration, as these users cover longer distances, and are better positioned to shop around to reduce the impact of local price fluctuations.

Probably the main upshot of our findings is that competition alone cannot reduce prices. Unless consumers engage with the market, the benefits of competition are less likely to be transferred to them. On the other hand, if they do engage with the market, then increasing market concentration is less likely to leave them facing increased margins even in concentrated markets. Moreover, engagement with the market is also important when concentration falls (as a result of entry). Although margins are likely to drop, they drop more in areas with more informed consumers.

These findings impact policy at two levels. First of all, the harm avoided by blocking a harmful increase in concentration includes avoiding some regressive distributional effects. Second, getting the market structure right may only offer a partial solution to a competition problem. Demand-side remedies may also be needed to ensure that all consumers engage with the market. Moreover, our findings also give support to arguments that even where blocking or breaking up concentration is not possible, demand-side remedies may help mitigate harmful effects, provided that some choice still exists for consumers.

8 Conclusion

Motor fuel is a non-trivial part of poorer households' expenditure, which means that poorer households already pay a larger share of their income on transport-related fuel. If they pay a higher price for increased market concentration, the impact is much more pronounced in relative terms. This is important because it implies that antitrust needs to revisit some of its conventional wisdom and account for the possibility that some people benefit more from the elimination of conduct that reduces competition, and this should be reflected in the design of remedies, i.e. remedies should not be designed with the average consumer in mind, but accounting for the heterogeneity of the impact of remedies across different income groups.

An important implication of our findings is that they offer support for the argument that antitrust could help address inequality while staying true to its mission of promoting competition.[16] We do not argue that income or wealth equality should be incorporated directly into competition policies. But we emphasise that ill-designed and executed competition policy and enforcement can contribute to increased inequality. Moreover, the success of the competition policy should not be evaluated for the average consumer. Instead, competition policy, when possible, should consider the possibility of a differential impact and impose remedies accordingly.

Motor fuel is similar to food in the sense that it is a non-discretionary part of household expenditure, which also displays significant distributional differences. Mattioli et al. (2018) identify a distinct group of households, around 10% of the UK's population who are in car-related economics stress: on low income, experience high motoring costs, and a low response to fuel price changes. This thinking is seemingly also gaining some consideration in the regulatory review of mergers. The UK Competition and Markets Authority specifically emphasised the difference in local areas regarding food and petrol expenditure (lower-income areas spending a relatively larger proportion of their income on food/petrol) in the Sainsbury's/Asda merger.[17] Whilst we think this is an important and welcome development, we also believe that more micro-level evidence is needed on this topic. To build up the evidentiary toolkit of competition authorities, we hope that this paper will help foster the drive to deliver more merger retrospectives that estimate not only the average but the distributional effects of mergers as well.

References

Allen, J., Clark, R. & Houde, J.-F. (2014), 'The effect of mergers in search markets: Evidence from the canadian mortgage industry', *American Economic Review* **104**(10), 3365–96.

Athey, S., Tibshirani, J., Wager, S. et al. (2019), 'Generalized random forests', *The Annals of Statistics* **47**(2), 1148–1178.

Athey, S. & Wager, S. (2019), 'Estimating treatment effects with causal forests: An application', *arXiv preprint arXiv:1902.07409* .

Baker, J. B. & Salop, S. C. (2015), 'Antitrust, competition policy, and inequality', *Geo. LJ Online* **104**, 1.

Borusyak, K. & Jaravel, X. (2017), 'Revisiting event study designs', *Available at SSRN 2826228* .

Breiman, L., Friedman, J., Stone, C. J. & Olshen, R. A. (1984), *Classification and regression trees*, CRC press.

Brown, J. R. & Goolsbee, A. (2002), 'Does the internet make markets more competitive? evidence from the life insurance industry', *Journal of political economy* **110**(3), 481–507.

Byrne, D. P. & de Roos, N. (2017), 'Consumer search in retail gasoline markets', *The Journal of Industrial Economics* **65**(1), 183–193.

Byrne, D. P. & Martin, L. A. (2021), 'Consumer search and income inequality', *International Journal of Industrial Organization* p. 102716.

Byrne, D. P., Nah, J. S. & Xue, P. (2018), 'Australia has the world's best petrol price data: Fuelwatch and fuelcheck', *Australian Economic Review* **51**(4), 564–577.

Caplovitz, D. (1963), The poor pay more: Consumer practices of low-income families, Technical report.

Cooper, T. E. & Jones, J. T. (2007), 'Asymmetric competition on commuter routes: the case of gasoline pricing', *Southern Economic Journal* pp. 483–504.

De los Santos, B. (2018), 'Consumer search on the internet', *International Journal of Industrial Organization* **58**, 66–105.

De los Santos, B., Hortaçsu, A. & Wildenbeest, M. R. (2012), 'Testing models of consumer search using data on web browsing and purchasing behavior', *American economic review* **102**(6), 2955–80.

[16] See for example Baker & Salop (2015), or Shapiro (2018).

[17] Para. 8.283

Diamond, P. (1987), 'Consumer differences and prices in a search model', *The Quarterly Journal of Economics* **102**(2), 429–436.

Dierx, A., Ilzkovitz, F., Pataracchia, B., Ratto, M., Thum-Thysen, A. & Varga, J. (2017), 'Does eu competition policy support inclusive growth?', *Journal of Competition Law & Economics* **13**(2), 225–260.

Ellison, G. & Ellison, S. F. (2009), 'Search, obfuscation, and price elasticities on the internet', *Econometrica* **77**(2), 427–452.

Ellison, G. & Fisher Ellison, S. (2005), 'Lessons about markets from the internet', *Journal of Economic perspectives* **19**(2), 139–158.

Ennis, S. F., Gonzaga, P. & Pike, C. (2019), 'Inequality: A hidden cost of market power', *Oxford Review of Economic Policy* **35**(3), 518–549.

Freyaldenhoven, S., Hansen, C. & Shapiro, J. M. (2019), 'Pre-event trends in the panel event-study design', *American Economic Review* **109**(9), 3307–38.

Hastings, J. S. (2004), 'Vertical relationships and competition in retail gasoline markets: Empirical evidence from contract changes in southern california', *American Economic Review* **94**(1), 317–328.

Hosken, D. S., McMillan, R. S. & Taylor, C. T. (2008), 'Retail gasoline pricing: What do we know?', *International Journal of Industrial Organization* **26**(6), 1425–1436.

Houde, J.-F. (2012), 'Spatial differentiation and vertical mergers in retail markets for gasoline', *American Economic Review* **102**(5), 2147–82.

Khan, L. M. & Vaheesan, S. (2017), 'Market power and inequality: The antitrust counterrevolution and its discontents', *Harv. L. & Pol'y Rev.* **11**, 235.

Kwoka, J. (2014), *Mergers, merger control, and remedies: A retrospective analysis of US policy*, MIT Press.

Lach, S. & Moraga-González, J. L. (2017), 'Asymmetric price effects of competition', *The Journal of Industrial Economics* **65**(4), 767–803.

Mariuzzo, F. & Ormosi, P. L. (2019), 'Post-merger price dynamics matters, so why do merger retrospectives ignore it?', *Review of Industrial Organization* **55**(3), 403–429.

Mattioli, G., Wadud, Z. & Lucas, K. (2018), 'Vulnerability to fuel price increases in the uk: A household level analysis', *Transportation Research Part A: Policy and Practice* **113**, 227–242.

Nishida, M. & Remer, M. (2018), 'The determinants and consequences of search cost heterogeneity: Evidence from local gasoline markets', *Journal of Marketing Research* **55**(3), 305–320.

Pennerstorfer, D., Schmidt-Dengler, P., Schutz, N., Weiss, C. & Yontcheva, B. (2020), 'Information and price dispersion: Theory and evidence', *International Economic Review* **61**(2), 871–899.

Roth, J. (2018), Pre-test with caution: Event-study estimates after testing for parallel trends, Technical report.

Rubin, D. B. (1974), 'Estimating causal effects of treatments in randomized and nonrandomized studies.', *Journal of educational Psychology* **66**(5), 688.

Salop, S. & Stiglitz, J. (1977), 'Bargains and ripoffs: A model of monopolistically competitive price dispersion', *The Review of Economic Studies* **44**(3), 493–510.

Schmidheiny, K. & Siegloch, S. (2019), 'On event study designs and distributed-lag models: Equivalence, generalization and practical implications'.

Sengupta, A. & Wiggins, S. N. (2014), 'Airline pricing, price dispersion, and ticket characteristics on and off the internet', *American Economic Journal: Economic Policy* **6**(1), 272–307.

Shapiro, C. (2018), 'Antitrust in a time of populism', *International Journal of Industrial Organization* **61**, 714–748.

Simpson, J. & Taylor, C. (2008), 'Do gasoline mergers affect consumer prices? the marathon ashland petroleum and ultramar diamond shamrock transaction', *The Journal of Law and Economics* **51**(1), 135–152.

Stahl, D. O. (1989), 'Oligopolistic pricing with sequential consumer search', *The American Economic Review* pp. 700–712.

Stango, V. & Zinman, J. (2016), 'Borrowing high versus borrowing higher: price dispersion and shopping behavior in the us credit card market', *The Review of Financial Studies* **29**(4), 979–1006.

Sun, L. & Abraham, S. (2020), 'Estimating dynamic treatment effects in event studies with heterogeneous treatment effects', *Journal of Econometrics* .

Tang, Z., Smith, M. D. & Montgomery, A. (2010), 'The impact of shopbot use on prices and price dispersion: Evidence from online book retailing', *International Journal of Industrial Organization* **28**(6), 579–590.

Taylor, C. T., Kreisle, N. M. & Zimmerman, P. R. (2010), 'Vertical relationships and competition in retail gasoline markets: Empirical evidence from contract changes in southern california: Comment', *American Economic Review* **100**(3), 1269–76.

Varian, H. R. (1980), 'A model of sales', *The American economic review* **70**(4), 651–659.

Wadud, Z., Graham, D. J. & Noland, R. B. (2010), 'Gasoline demand with heterogeneity in household responses', *The Energy Journal* **31**(1).

Wager, S. & Athey, S. (2018), 'Estimation and inference of heterogeneous treatment effects using random forests', *Journal of the American Statistical Association* **113**(523), 1228–1242.

Wilson, C. M. & Price, C. W. (2010), 'Do consumers switch to the best supplier?', *Oxford Economic Papers* **62**(4), 647–668.

Woodward, S. E. & Hall, R. E. (2012), 'Diagnosing consumer confusion and sub-optimal shopping effort: Theory and mortgage-market evidence', *American Economic Review* **102**(7), 3249–76.

Zac, A., Casti, C., Decker, C. & Ezrachi, A. (2020), 'Competition law and income inequality: A panel data econometric approach', *Available at SSRN* .

Appendix

A Simulations to demonstrate our ensemble bagging model of causal forests

We conducted some experiments to justify using an ensemble causal forest method in cases where the data has a relatively small number of individuals and many potential sources of treatment heterogeneity. We simulate a dataset with $n = 1000$ individuals, where the number of factors increases from $k = 20$ to $k = 120$. In each loop, the treatment effect is a linear function of $J = k/5$ of these factors.

Our data generating process (DGP) is as follows:

$$Y_i = \alpha X_i + \beta_W W_i + \beta W_i X_i + U_i \tag{11}$$

Where W_i is the treatment variable, following a binomial distribution $W_i \sim B(n, 0.5)$, and $U_i \sim N(0, 1)$. The variable X_i represents a vector of J covariates, generated from a multivariate normal distribution. β is a vector of J parameters with $\beta_j = 1$ (for $j = 1, \ldots, J$) (i.e. as we increase J we add add covariates. The starting DGP is defined as $J = 2$; $k = 10$, $N = 1000$.

Assume that we are interested in the effect of a set of factors X_1, X_2, X_3, where theory supports some relationship between the factors and the treatment effect. For this reason we then record estimates of $\beta_{1,2,3}$, as we systematically vary the k (20,30,40,50,60,70,80,90,100), and correspondingly the J (the number of variables causing heterogeneity in treatment) parameter (4,6,8,10,12,14,16,18,20), while holding everything else constant.

We compare the following two processes:

- **Single causal forest:** We follow Athey & Wager (2019) and start by training two random forests for Y and W and use its parameters as parameter choices to run our causal forest. Similarly to Athey & Wager (2019) we first train a pilot causal forest including all features, and then train a second forest only on those features that had most splits in the first forest (features that had at least the average share of splits). This helps in focusing efforts on the most important features. Our change in comparison to this formula is that we force our feature of interest ($X_{1,2,3}$) to be in the second, smaller pool of features as we are specifically interested in their role in treatment heterogeneity.
- **Ensemble causal forests:** We estimate $C = 1000$ causal forests. In each iteration, we repeatedly draw a random sample of $J/10$ features, plus we add our features of interest and estimate the causal forest on this small sample of features. The idea is that through our iterations, each feature has interacted with $X_{1,2,3}$. We then average over the estimates to give our ensemble estimate. For example for X_1 we get $\hat{\beta}_1 = 1/C \sum_{c=1}^{C} \hat{\beta_{c,1}}$.

Figure 8 shows the estimates (and standard errors) for $\hat{\beta}_1$. The horizontal red line marks the true effect β_1. Using the single forest method, the estimates drop as we have an increasing number of features and a small sample size. Using the ensemble method, the estimates are not affected by the increase in the number of features.

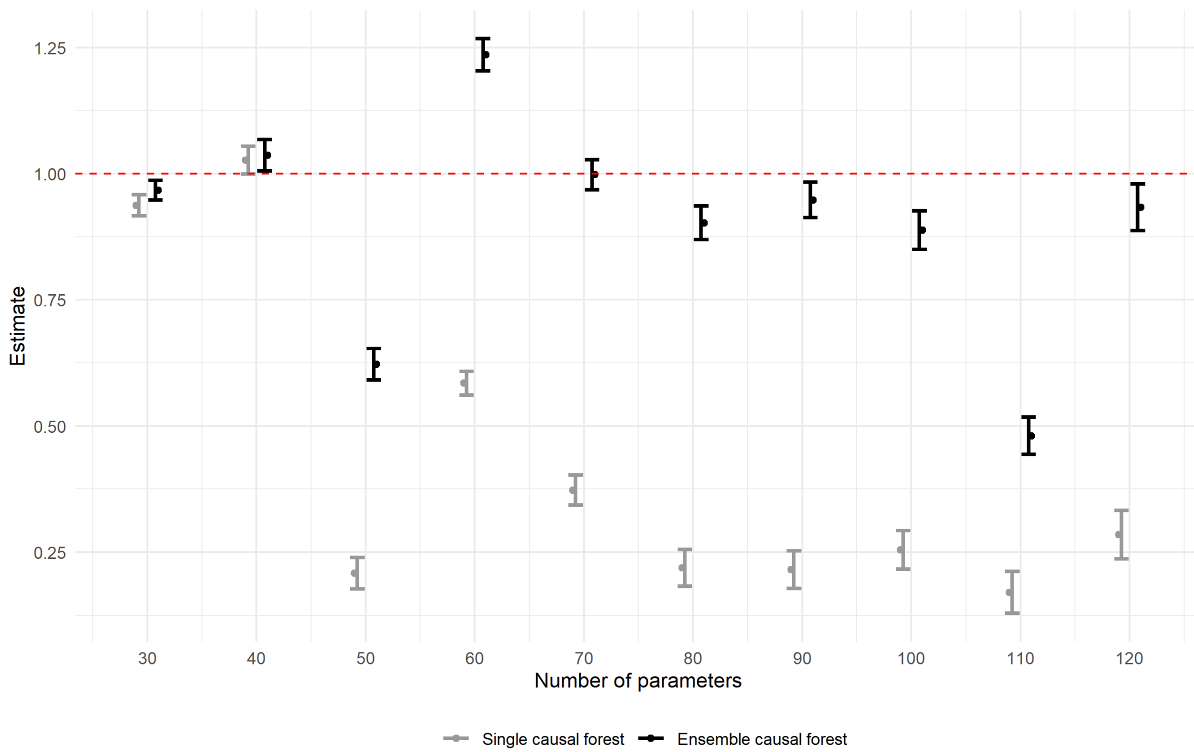

Figure 8: Ensemble v single causal forest - simulation results

B Figures and tables for descriptive part

B.1 Tables

Table B.1: Main features of the data

topic	variable	topic	variable
time	year	**housing**	average rent
	quarter		weekly rent $1-74
wealth	gini coefficient		weekly rent $75-99
	income share of top 1%		weekly rent $100-124
	income share of top 5%		weekly rent $125-149
	income share of top 10%		weekly rent $150-174
	% of people in lowest quartile (relative to AUS)		weekly rent $175-199
	% of people in second quartile (relative to AUS)		weekly rent $200-224
	% of people in third quartile (relative to AUS)		weekly rent $225-249
	% of people in highest quartile (relative to AUS)		weekly rent $250-274
employment	number of employed people living in region no		weekly rent $275-299
	number of earners		weekly rent $300-324
	age of earners		weekly rent $325-349
	sum income		weekly rent $350-374
	median income		weekly rent $375-399
	mean income		weekly rent $400-424
business income/cost	mean ind income before tax		weekly rent $425-449
	mean total business income		weekly rent $450-549
	mean total business expense		weekly rent $550-649
	mean net business income		weekly rent $650-749
	mean estimated business tax		weekly rent $750-849
	mean gross rent		weekly rent $850-949
	mean net rent		weekly rent $950 and over
age	age 0-4		nil rent payments
	age 4-10	**internet access**	internet accessed from dwelling (%)
	age 10-15		internet not accessed from dwelling (%)
	age 15-20		internet access from home / population
	age 20-25	**number of cars**	no cars
	age 25-30		one motor vehicle
	age 30-35		two motor vehicles
	age 35-40		three motor vehicles
	age 40-45		four or more
	age 45-50		average no cars
	age 50-55	**commuting**	average commuting distance (mi)
	age 55-60		median commuting distance (mi)
	age 60-65		interquartile range of commuting (mi)
	age 65-70		standard deviation of commuting (mi)
	age 70-75		train
	age 75-80		bus
	age 80-85		ferry
	age 85-99		tram
	age 65+		taxi
	age 35-65		car as driver
	age 15-35		car as passenger
	age 0-15		truck
	population		motorbike scooter
education	index of education and occupation		bicycle
	advanced diploma and diploma level		walked only
	bachelor degree level		worked at home
	certificate I II level		did not go to work
	certificate III IV level	**competition**	number of rivals within 1 mile
	certificate level		number of rivals within 2 mile
	certificate level nfd		number of rivals within 5 mile
	graduate diploma and graduate certificate level	**brand**	brand size
	level of education not stated		top brand (bp, shell, caltex)
	postgraduate degree level		number of same brand stations within 1 mile
	index of economic resources		number of same brand stations within 2 mile
	index of relative socio-economic advantage and disadvantage		number of same brand stations within 5 mile
	index of relative socio-economic disadvantage		

Table B.2: Summary statistics of the main variables

	mean	**sd**	**10%**	**25%**	**50%**	**75%**	**90%**
ulp retail price	138.28	10.30	127.29	131.92	137.16	143.72	151.05
diesel retail price	146.19	8.80	136.49	140.77	145.46	150.56	156.83
ulp wholesale price	123.54	6.85	115.94	119.65	123.61	127.78	131.78
diesel wholesale price	129.62	6.77	122.01	125.62	129.46	133.69	137.08
cushing price	25.81	2.75	22.85	24.55	25.84	27.07	28.53
ulp margin	1.12	0.07	1.06	1.08	1.10	1.15	1.21
diesel margin	1.13	0.06	1.07	1.09	1.12	1.15	1.19
number of rivals (within 1mi)	1.85	1.70	0	1	2	3	4
number of rivals (within 2mi)	5.19	4.24	0	2	5	8	10
number of rivals (within 5mi)	21.82	20.94	1	4	14	37	56
median income	50607.22	9083.98	41979	45225	50049	54605	60254
mean income	64962.14	18375.49	51568	55371	61552	68390	79770
usual resident population	12249.66	7150.07	4297	5870	11790	16517	23065
people aged 0-14 years	18.90	4.61	14.3	16.7	19.2	21.7	24.4
people aged 15-64 years	66.35	9.08	60.7	63.2	67	70.4	73.6
people aged 65 years and over	14.75	8.85	6.2	10.1	14.1	18.4	21.1
median age	38.46	6.57	32.2	33.5	37.6	41.5	44.7
sex ratio	109.50	41.00	92	96.9	99.7	105.1	120
earners age	43.07	4.63	37	39	44	47	48
number of earners	7123.57	4754.28	2040	3303	6589	9942	13621
no educational attainment	53.46	72.83	3	10	26	67	136
average commuting distance (mi)	11.27	8.34	4.33	6.18	9.36	13.95	18.33
median commuting distance (mi)	6.74	4.65	1.79	2.79	6.55	8.96	12.99
car as driver	0.28	0.08	0.22	0.25	0.29	0.32	0.34
one motor vehicle	0.25	0.07	0.17	0.20	0.26	0.30	0.32
index of relative socio-economic disadvantage	993.10	80.00	917	975	997	1040	1071
index of relative socio-economic advantage and disadvantage	991.43	74.62	901	956	988	1041	1084
index of economic resources	1003.85	79.78	925	973	1016	1050	1089
index of education and occupation	980.20	73.46	886	928	977	1016	1101

The price statistics are reported for 489,721 weekly observations, the area characteristics are reported for the 1053 distinct petrol station locations.

Table B.3: Difference in ULP and diesel margins by the main variables of interest

	petrol		**diesel**	
commuting distance	**low**	**high**	**low**	**high**
	1.14	1.104	1.144	1.116
	(0.082)	(0.056)	(0.063)	(0.045)
competition 1mi	low	high	low	high
	1.121	1.118	1.129	1.127
	(0.074)	(0.065)	(0.057)	(0.052)
competition 2mi	low	high	low	high
	1.136	1.099	1.136	1.119
	(0.08)	(0.051)	(0.063)	(0.042)
competition 5mi	low	high	low	high
	1.149	1.088	1.142	1.113
	(0.082)	(0.037)	(0.065)	(0.037)
income	low	high	low	high
	1.133	1.11	1.134	1.126
	(0.072)	(0.07)	(0.058)	(0.054)
education	low	high	low	high
	1.135	1.109	1.135	1.126
	(0.078)	(0.064)	(0.062)	(0.05)
% of people +65 age	low	high	low	high
	1.129	1.114	1.137	1.123
	(0.086)	(0.055)	(0.065)	(0.046)
% people internet home	low	high	low	high
	1.133	1.111	1.135	1.124
	(0.082)	(0.061)	(0.061)	(0.051)

Standard deviation in parentheses.

Table B.4: Comparison of treatment and control

	control	treatment
income	49298.52	50032.73
	(9165.062)	(8772.878)
earners age	43.392	43.392
	(4.662)	(4.519)
number of earners	7327.915	7465.363
	(4944.795)	(4452.06)
Median commuting distance (kms)	10.806	9.57
	(7.356)	(6.946)
gini coefficient	0.471	0.473
	(0.053)	(0.057)
Car as driver	0.276	0.291
	(0.048)	(0.042)
One motor vehicle	0.259	0.265
	(0.059)	(0.055)
Usual Resident Population	12251.86	12336.72
	(7327.96)	(6777.509)
Index of Education and Occupation	976.245	983.53
	(88.463)	(80.113)
competitors (1mi)	2.405	2.558
	(2.249)	(1.805)
competitors (2mi)	5.392	7.128
	(4.617)	(4.343)
competitors (5mi)	20.772	23.543
	(22.499)	(20.502)

Standard deviation in parentheses.

Table B.5: Number of exits and entries by brand

brand	exitter	exitter_market	entrants	entrant_market	total number of petrol stations in the sample
BP	92	42	90	40	260
Caltex	119	49	84	38	254
Shell	111	39	47	23	163
Puma	10	4	46	20	89
Independent	38	15	30	15	75
Gull	68	25	10	3	61
Caltex Woolworths	2	1	38	15	57
Coles Express	4	1	24	11	53
7-Eleven	1	1	37	16	46
Ampol	54	17	7	3	36
Mobil	40	11	14	5	35
Liberty	39	17	11	8	31
Peak	21	9	9	6	30
United	6	1	7	5	26
Vibe	0	0	18	6	20
Kleenheat	29	9	0	0	13
Wesco	20	7	0	0	12
Better Choice	0	0	2	1	8
Amgas	8	6	1	1	6
Eagle	4	2	7	2	6
BOC	4	3	4	1	5
Kwikfuel	4	2	3	1	5
Oasis	0	0	1	1	2
Swan Taxis	4	2	0	0	2
Black and White	0	0	0	0	1
FastFuel 24/7	3	1	3	1	1
Metro Petroleum	0	0	2	1	1
United Fuels West	0	0	0	0	1
TOTAL	681	264	495	223	1299

Table B.6: Margin by brand

brand	frequency	petrol	petrol_margin	diesel	diesel_margin
Eagle	6	139.642	1.164	142.211	1.156
Independent	75	133.078	1.147	137.958	1.147
United Fuels West	1	129.847	1.123	132.161	1.119
BP	260	129.290	1.114	133.051	1.131
Shell	163	127.676	1.108	132.955	1.129
Coles Express	53	136.030	1.103	143.913	1.144
Ampol	36	121.979	1.097	127.092	1.118
Caltex	254	127.026	1.095	131.814	1.121
7-Eleven	46	137.068	1.090	126.891	1.116
Caltex Woolworths	57	131.436	1.089	134.804	1.114
Gull	61	122.681	1.086	127.297	1.105
Wesco	12	115.058	1.086	119.853	1.104
United	26	129.351	1.081	133.313	1.098
Liberty	31	122.115	1.079	126.795	1.092
Vibe	20	130.425	1.078	134.635	1.091
Puma	89	126.971	1.072	132.927	1.106
Kwikfuel	5	121.896	1.058	127.323	1.089
Metro Petroleum	1	135.732	1.056	146.029	1.085
Peak	30	121.682	1.056	127.388	1.085
Better Choice	8	126.182	1.050	131.054	1.065
Amgas	6	114.070	1.048	120.848	1.088
Mobil	35	120.515	1.045	127.508	1.080
FastFuel 24/7	1	124.164	1.042	128.275	1.064
Oasis	2	118.343	1.037	124.152	1.070

Table B.7: Margin by competition

level of competition within				
1 mile	**2 miles**	**5 miles**	**ulp margin**	**diesel margin**
low	low	low	1.150	1.141
high	low	low	1.168	1.153
low	high	low	1.130	1.129
high	high	low	1.147	1.137
low	low	high	1.087	1.111
high	low	high	1.086	1.105
low	high	high	1.085	1.113
high	high	high	1.091	1.115

We split the number of competitors within each radius around their median values. For within 1 mile: 0-2 (low) versus 3 or more (high) competitors, for within 2 miles: 0-5 (low) versus 6 or more (high) competitors, and for within 5 miles: 0-15 (low) versus 16 or more (high) competitors.

Table B.8: Margin by brand by competition by income

station	low income	med income	high income	station	low income	med income	high income
7-Eleven	1.088	1.089	1.084	**Liberty**	1.111	1.101	1.124
n	6	20	9	n	9	11	8
comp 1mi	1.761	1.899	2.764	comp 1mi	1.022	2.389	2.261
comp 5mi	46.821	33.118	31.138	comp 5mi	13.911	22.951	46.469
Ampol	1.158	1.076	1.091	**Mobil**	1.081	1.072	1.068
n	15	6	13	n	5	15	7
comp 1mi	1.683	1.866	2.515	comp 1mi	2.158	2.584	3.149
comp 5mi	7.092	53.11	42.823	comp 5mi	47.131	29.382	47.531
BP	1.169	1.099	1.136	**Peak**	1.118	1.079	1.072
n	66	61	70	n	4	11	13
comp 1mi	2.476	1.548	2.215	comp 1mi	1.976	1.069	1.252
comp 5mi	13.206	28.724	28.633	comp 5mi	30.032	18.292	24.896
Caltex	1.14	1.099	1.12	**Puma**	1.11	1.078	1.091
n	71	59	64	n	15	26	24
comp 1mi	2.254	1.646	2	comp 1mi	1.269	1.854	2.687
comp 5mi	12.741	30.465	26.431	comp 5mi	19.933	28.445	29.227
Caltex Woolworths	1.102	1.081	1.123	**Shell**	1.169	1.1	1.133
n	12	13	13	n	51	33	45
comp 1mi	2.491	0.92	2.195	comp 1mi	2.129	1.434	1.659
comp 5mi	21.877	27.859	18.925	comp 5mi	11.983	19.552	27.305
Coles Express	1.108	1.09	1.112	**United**	1.118	1.064	1.094
n	15	15	18	n	7	7	6
comp 1mi	2.356	2.203	2.431	comp 1mi	4.01	0.805	1.663
comp 5mi	31.78	43.328	32.121	comp 5mi	11.671	28.136	18.224
Gull	1.13	1.108	1.106	**Vibe**	1.109	1.084	1.054
n	23	14	14	n	6	8	5
comp 1mi	2.653	0.946	1.882	comp 1mi	1.756	1.012	3.53
comp 5mi	11.36	17.385	41.297	comp 5mi	8.382	8.85	43.091
Independent	1.184	1.144	1.206				
n	31	17	12				
comp 1mi	1.199	0.752	2.365				
comp 5mi	3.491	3.343	18.521				

Table B.9: Mean margin by income, competition, and population size

		low competition		high competition	
		low population	high population	low population	high population
ulp	low income	1.149	1.102	1.151	1.121
		(0.081)	(0.05)	(0.074)	(0.051)
	high income	1.135	1.099	1.096	1.111
		(0.085)	(0.061)	(0.05)	(0.069)
diesel	low income	1.144	1.115	1.142	1.122
		(0.066)	(0.038)	(0.065)	(0.04)
	high income	1.133	1.123	1.117	1.13
		(0.062)	(0.05)	(0.043)	(0.052)

Standard deviation in parentheses.

B.2 Figures

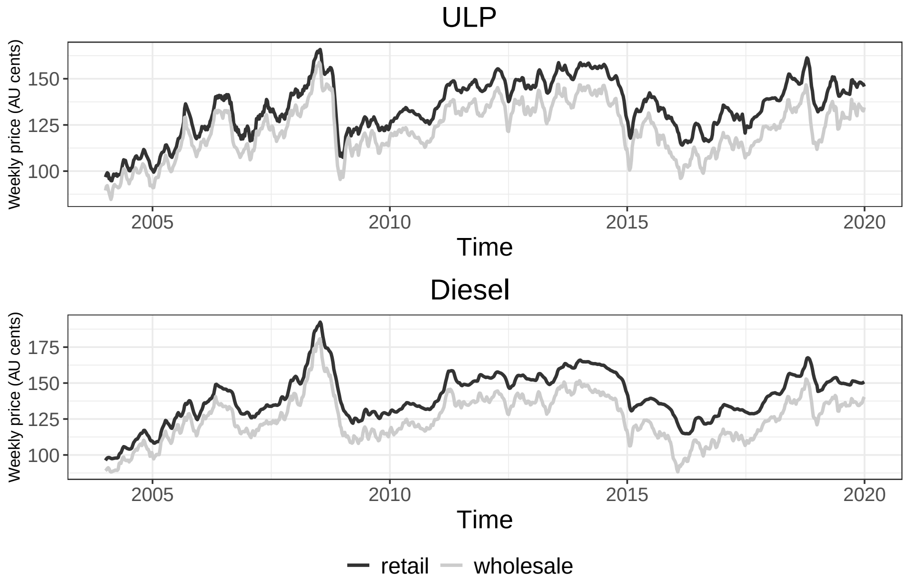

Figure B.9: Weekly retail ULP and diesel retail and wholesale prices over time

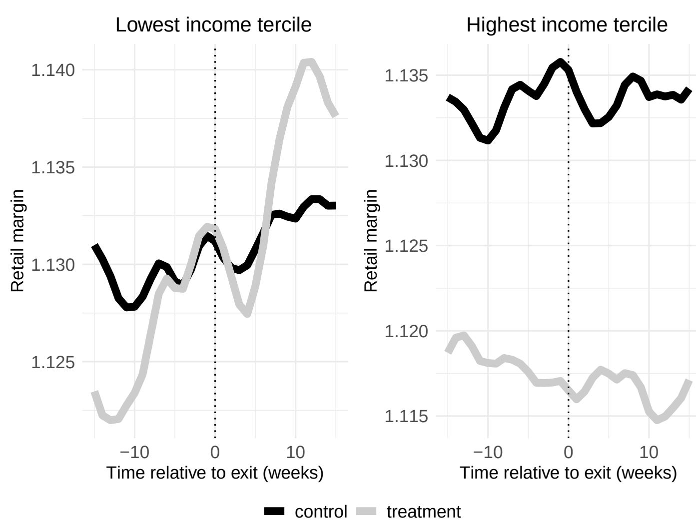

Figure B.10: Retail price margin before and after exit for diesel at different income levels

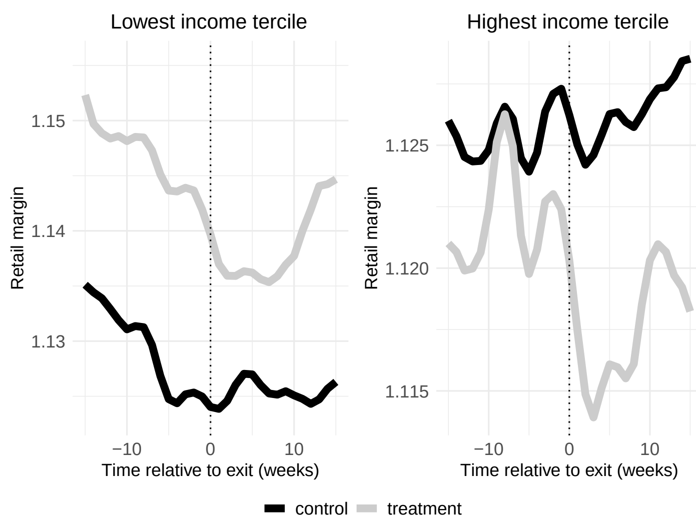

Figure B.11: Retail price margin before and after entry for diesel at different income levels

C Tables and figures for the Results section

C.1 Tables

Table C.1: Number of rivals in the exit and entry samples used for the causal forest estimation

	exit	**entry**
number of rivals within 1 mile	2.862	2.579
	(2.216)	(1.789)
number of rivals within 2 mile	6.605	5.690
	(4.066)	(5.517)
number of rivals within 5 miles	23.268	15.241
	(19.893)	(21.687)

Table C.2: Top 30 most important features for ULP and Diesel exit

variable	**ulp exit**	**variable**	**diesel exit**
Ferry	0 146	three motor vehicles	0 157
Certificate I II Level	0 13	four or more	0 15
mean estimated business tax	0 097	Ferry	0 144
average no cars	0 089	two motor vehicles	0 121
three motor vehicles	0 081	age 59	0 087
four or more	0 081	weekly residential rent 325-349	0 084
age 35	0 076	Certificate III IV Level	0 08
Motorbike scooter	0 074	one motor vehicle	0 078
age 65	0 065	Motorbike scooter	0 076
Did not go to work	0 06	Internet not accessed from dwelling	0 073
Standard deviation kms	0 059	age 70	0 071
median income	0 059	Certificate I II Level	0 07
age 60	0 058	Did not go to work	0 07
age 40	0 057	average no cars	0 069
age 70	0 057	Certificate Level	0 068
Bus	0 057	Car as driver	0 065
two motor vehicles	0 057	age 35-65	0 061
mean net business income	0 056	median income	0 06
age 75	0 055	age 40	0 058
age 80	0 05	mean net business income	0 057
age 85-99	0 047	age 45	0 057
age 20	0 047	weekly residential rent 125-149	0 057
Certificate III IV Level	0 047	age 85-99	0 056
IEO	0 047	Walked only	0 056
age 15	0 043	weekly residential rent 100-124	0 053
Car as driver	0 043	weekly residential rent 350-374	0 053
age 65 PLUS	0 041	mean estimated business tax	0 052
Certificate Level	0 041	comp5mi	0 052
weekly residential rent 125-149	0 041	age 04	0 051
Median commuting distance kms	0 04	age 35	0 051

The figures refer to importance score of each variable. It is calculated as the decrease in node impurity weighted by the probability of reaching that node.

Table C.3: Top 30 most important features for ULP and Diesel entry

variable	ulp exit	variable	diesel exit
Median commuting distance kms	0.151	age 70	0 136
age 65+	0.144	one motor vehicle	0 115
Interquartile range kms	0.143	three motor vehicles	0 113
Advanced Diploma and Diploma Level	0.143	four or more	0 11
age 75	0.141	age 75	0 103
age verage commuting distance kms	0.138	age 65+	0 103
Car as driver	0.133	mean net rent	0 097
Walked only	0.112	Bus	0 097
age 80	0.109	age 59	0 091
mean-net-rent	0.103	two motor vehicles	0 089
Index of Economic Resources	0.097	age 04	0 086
age 40	0.09	age 80	0 081
Car as passenger	0.09	Certificate I II Level	0 08
Motorbike scooter	0.083	Advanced Diploma and Diploma Level	0 079
Index of Relative Socio-econ Adv and Disadv	0.08	Index of Economic Resources	0 076
age 85-99	0.079	Nil payments	0 068
Ferry	0.079	age 65	0 067
age 65	0.078	age 10	0 066
average-no-cars	0.078	age 45	0 066
age 70	0.077	age 25	0 064
age 35	0.075	None	0 063
age 60	0.072	Interquartile range kms	0 062
Certificate I II Level	0.071	age 60	0 062
four or more	0.071	Truck	0 062
Index of Relative Socio-econ Disadv	0.066	age verage commuting distance kms	0 061
Bus	0.066	age 40	0 061
age 45	0.057	age 0-15	0 06
age 15	0.056	weekly residential rent 250-274	0 058
age 50	0.056	weekly residential rent 325-349	0 058
age 0-15	0.055	weekly residential rent 100-124	0 056
age 59	0.053	weekly residential rent 225-249	0 056

The figures refer to importance score of each variable. It is calculated as the decrease in node impurity weighted by the probability of reaching that node.

Table C.4: Predicted treatment effects of exit in Diesel, by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.222	0.173	0.194	0.143
		(0.114)	(0.098)	(0.102)	(0.085)
	high income	0.161	0.11	0.152	0.1
		(0.089)	(0.075)	(0.081)	(0.066)
high competition	**low income**	0.093	0.068	0.091	0.06
		(0.079)	(0.066)	(0.074)	(0.061)
	high income	0.056	0.026	0.068	0.033
		(0.067)	(0.056)	(0.063)	(0.052)

Bootstrapped standard errors in parentheses.

Table C.5: Predicted treatment effects of entry in Diesel, by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	-0.098	-0.089	-0.102	-0.09
		(0.044)	(0.038)	(0.049)	(0.044)
	high income	-0.095	-0.091	-0.083	-0.075
		(0.043)	(0.037)	(0.049)	(0.045)
high competition	**low income**	-0.08	-0.073	-0.076	-0.065
		(0.038)	(0.034)	(0.044)	(0.04)
	high income	-0.079	-0.076	-0.058	-0.052
		(0.038)	(0.033)	(0.045)	(0.042)

Bootstrapped standard errors in parentheses.

Table C.6: Predicted treatment effects of exit in ULP, by different levels of competition, income, commute and % of people age 65+

		low % of age 65+		high % of age 65+	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.246 (0.124)	0.205 (0.114)	0.275 (0.131)	0.249 (0.125)
	high income	0.108 (0.096)	0.068 (0.084)	0.148 (0.093)	0.122 (0.087)
high competition	**low income**	0.116 (0.079)	0.09 (0.069)	0.121 (0.082)	0.108 (0.076)
	high income	0.028 (0.065)	0.003 (0.057)	0.047 (0.058)	0.034 (0.052)

Bootstrapped standard errors in parentheses.

Table C.7: Predicted treatment effects of exit in ULP, by different levels of competition, income, commute and education

		low education		high education	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.318 (0.143)	0.285 (0.136)	0.215 (0.098)	0.191 (0.093)
	high income	0.171 (0.117)	0.138 (0.11)	0.093 (0.081)	0.07 (0.074)
high competition	**low income**	0.175 (0.1)	0.157 (0.093)	0.09 (0.064)	0.08 (0.059)
	high income	0.08 (0.083)	0.062 (0.075)	0.015 (0.054)	0.006 (0.05)

Bootstrapped standard errors in parentheses.

Table C.8: Predicted treatment effects of exit in ULP, by different levels of competition, income, and search - placebo treatment

		low internet		high internet	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.222 (0.114)	0.173 (0.098)	0.194 (0.102)	0.143 (0.085)
	high income	0.161 (0.089)	0.11 (0.075)	0.152 (0.081)	0.1 (0.066)
high competition	**low income**	0.093 (0.079)	0.068 (0.066)	0.091 (0.074)	0.06 (0.061)
	high income	0.056 (0.067)	0.026 (0.056)	0.068 (0.063)	0.033 (0.052)

Bootstrapped standard errors in parentheses.

Table C.9: Predicted treatment effects of **exit** in **ULP**, by different levels of competition, income, and search - [-5,+15] event window

		low internet		high internet	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.291 (0.099)	0.255 (0.091)	0.262 (0.099)	0.239 (0.089)
	high income	0.121 (0.089)	0.078 (0.085)	0.094 (0.09)	0.069 (0.085)
high competition	**low income**	0.161 (0.072)	0.152 (0.065)	0.149 (0.075)	0.15 (0.065)
	high income	0.061 (0.061)	0.044 (0.058)	0.048 (0.064)	0.045 (0.06)

Bootstrapped standard errors in parentheses.

Table C.10: Predicted treatment effects of **entry** in **ULP**, by different levels of competition, income, and search - [-5,+15] event window

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	-0.287 (0.068)	-0.313 (0.059)	-0.348 (0.066)	-0.366 (0.058)
	high income	-0.282 (0.062)	-0.314 (0.056)	-0.33 (0.059)	-0.355 (0.054)
high competition	**low income**	-0.211 (0.042)	-0.241 (0.037)	-0.272 (0.045)	-0.292 (0.04)
	high income	-0.208 (0.036)	-0.245 (0.035)	-0.256 (0.04)	-0.284 (0.037)

Bootstrapped standard errors in parentheses.

Table C.11: Predicted treatment effects of exit (within 2 miles) in ULP, by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.194 (0.131)	0.162 (0.111)	0.17 (0.12)	0.15 (0.097)
	high income	0.101 (0.099)	0.082 (0.077)	0.077 (0.1)	0.068 (0.074)
high competition	**low income**	0.145 (0.091)	0.107 (0.078)	0.139 (0.081)	0.112 (0.067)
	high income	0.078 (0.055)	0.047 (0.046)	0.069 (0.055)	0.046 (0.044)

Bootstrapped standard errors in parentheses.

Table C.12: Predicted treatment effects of exit (within 5 miles) in ULP, by different levels of competition, income, and search

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.123 (0.088)	0.085 (0.064)	0.053 (0.086)	0.025 (0.06)
	high income	0.068 (0.08)	0.05 (0.059)	0.018 (0.086)	0.005 (0.06)
high competition	**low income**	0.067 (0.085)	0.062 (0.064)	0.011 (0.079)	0.011 (0.058)
	high income	0.009 (0.079)	0.029 (0.06)	-0.029 (0.081)	-0.009 (0.059)

Bootstrapped standard errors in parentheses.

Table C.13: Predicted treatment effects of **exit** in **ULP**, by different levels of competition, income, and search - clustered by postcode

		low internet		**high internet**	
		low commute	**high commute**	**low commute**	**high commute**
low competition	**low income**	0.168 (0.076)	0.154 (0.066)	0.162 (0.086)	0.159 (0.076)
	high income	0.105 (0.072)	0.086 (0.063)	0.082 (0.083)	0.074 (0.072)
high competition	**low income**	0.105 (0.08)	0.078 (0.053)	0.087 (0.073)	0.073 (0.054)
	high income	0.055 (0.076)	0.025 (0.049)	0.021 (0.071)	0.004 (0.051)

Table C.14: Predicted treatment effects of **entry** in **ULP**, by different levels of competition, income, and search - clustered by postcode

		low internet		high internet	
		low commute	high commute	low commute	high commute
low competition	low income	-0.332 (0.05)	-0.325 (0.045)	-0.37 (0.056)	-0.36 (0.051)
	high income	-0.333 (0.04)	-0.329 (0.036)	-0.366 (0.042)	-0.358 (0.038)
high competition	low income	-0.286 (0.044)	-0.279 (0.04)	-0.326 (0.048)	-0.317 (0.044)
	high income	-0.292 (0.033)	-0.288 (0.03)	-0.327 (0.036)	-0.319 (0.032)

Table C.15: Predicted treatment effects of **exit** in **ULP**, by different levels of competition, income, and search - nearest 2 control firms

		low internet		high internet	
		low commute	high commute	low commute	high commute
low competition	low income	0.004 (0.061)	-0.022 (0.033)	0.024 (0.057)	0.007 (0.034)
	high income	0.151 (0.091)	0.125 (0.081)	0.168 (0.079)	0.15 (0.068)
high competition	low income	0.134 (0.082)	0.08 (0.06)	0.171 (0.076)	0.128 (0.056)
	high income	0.337 (0.136)	0.282 (0.123)	0.367 (0.12)	0.322 (0.106)

Table C.16: Predicted treatment effects of **entry** in **ULP**, by different levels of competition, income, and search - nearest 2 control firms

		low internet		high internet	
		low commute	high commute	low commute	high commute
low competition	low income	-0.302 (0.06)	-0.332 (0.059)	-0.374 (0.058)	-0.388 (0.056)
	high income	-0.281 (0.053)	-0.319 (0.055)	-0.346 (0.053)	-0.367 (0.052)
high competition	low income	-0.237 (0.041)	-0.271 (0.043)	-0.313 (0.045)	-0.33 (0.044)
	high income	-0.22 (0.037)	-0.262 (0.041)	-0.288 (0.042)	-0.312 (0.041)

C.2 Figures

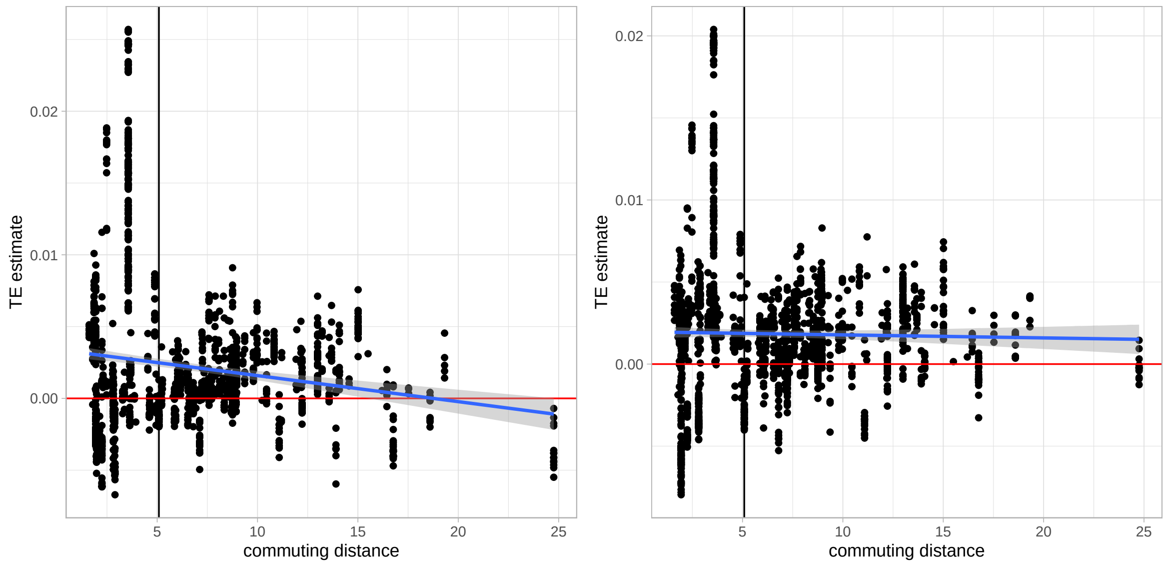

Figure C.12: Treatment effects by commuting distance

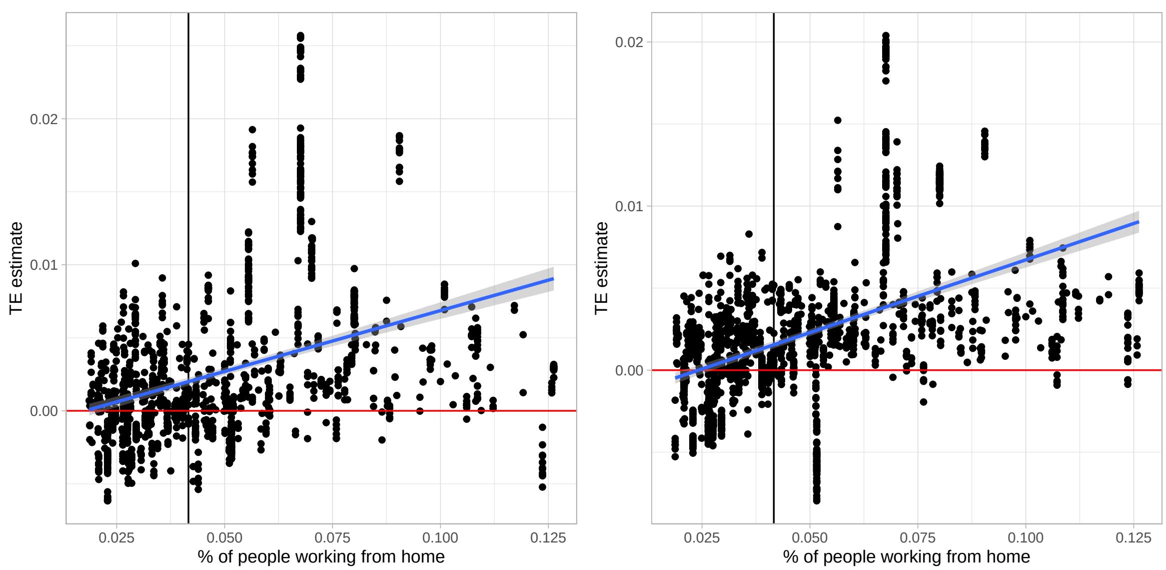

Figure C.13: Treatment effects by home working prevalence

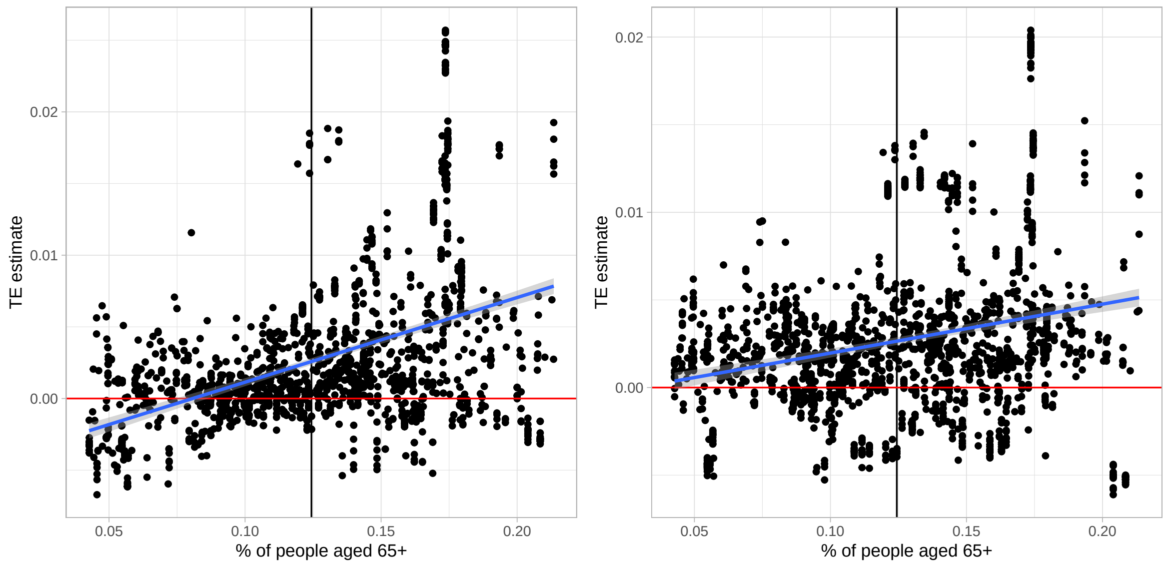

Figure C.14: Treatment effects by % of people aged 65+

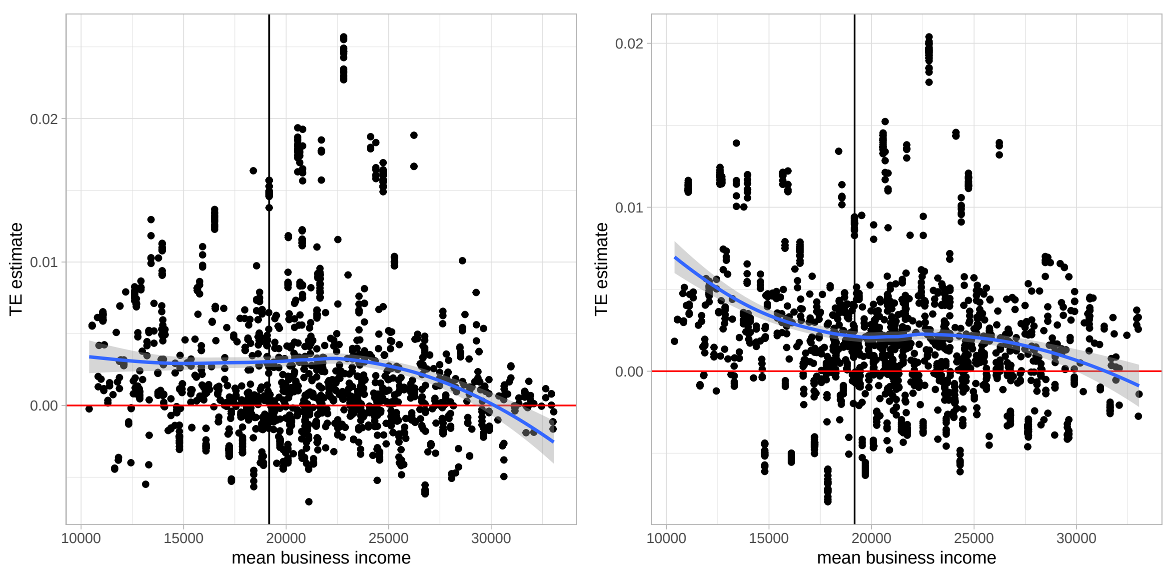

Figure C.15: Treatment effects by business income

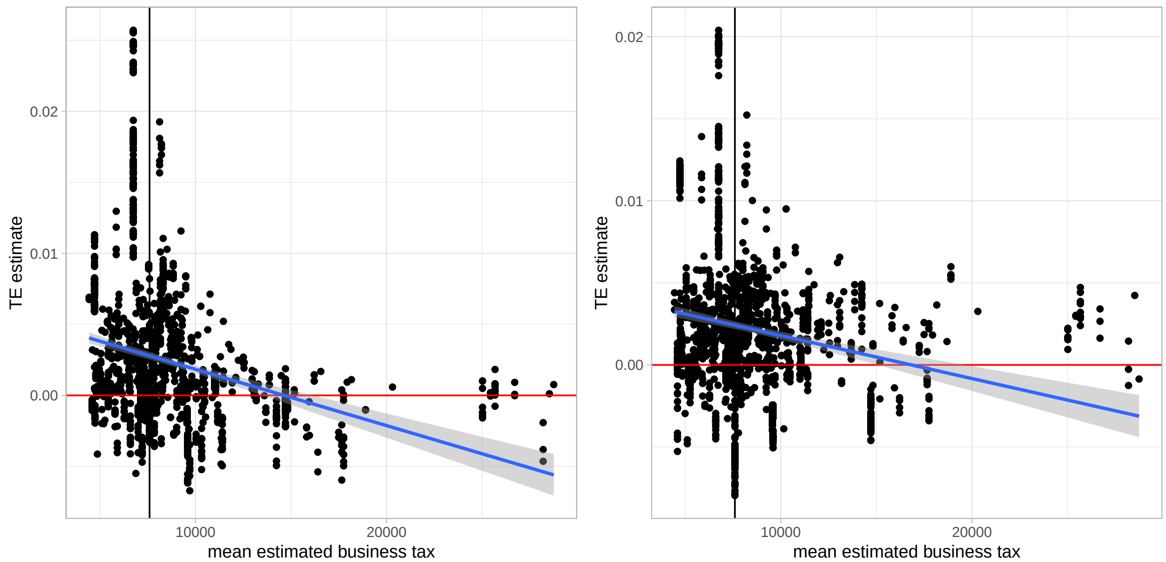

Figure C.16: Treatment effects by business tax

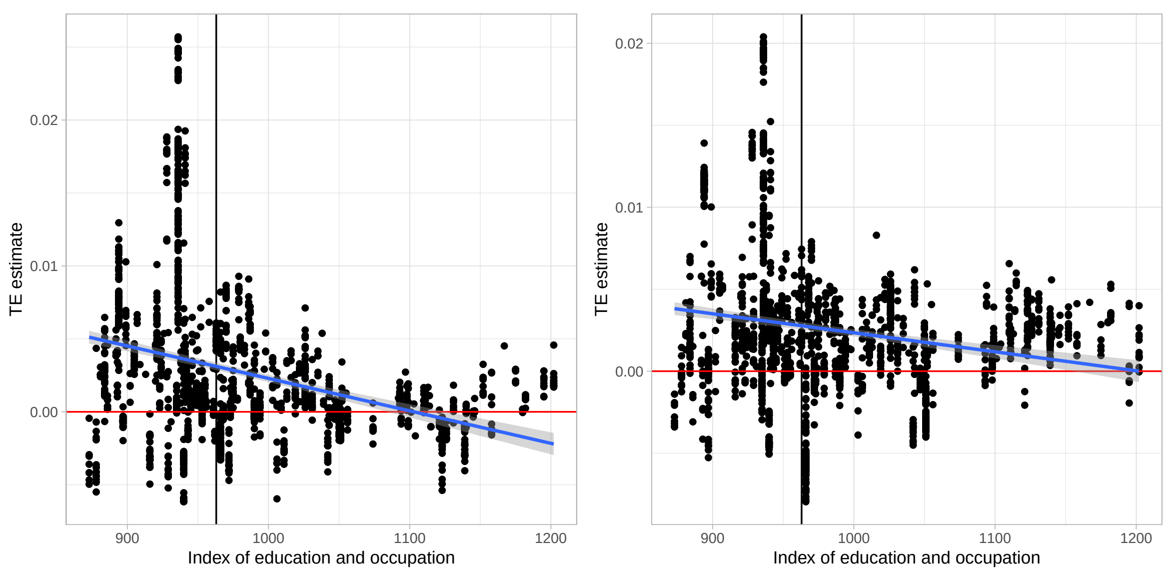

Figure C.17: Treatment effects by educational level